

Tutorial 3

Distributed Transactions Management

+ Q1 Distributed Transaction

1

Olympic Game Database **Event**(EventID, ***)
 Athlete(CompID, ***)

1. **Results**(EventID, CompID, Position)
2. **Medals**(EventID, Medal, CompID) //Medal: Gold/Silver/Bronze
3. **MedalTally**(Country, Medal, Number)
4. **Competitors**(CompID, Country)
5. **Medalists**(CompID, NMedals) // NMedals: total of medals won

- When a race is running, a number of tables need to be updated.
 - Assume that when an event is completed, a file is created at the venue giving CompID and Position associated with that EventID, then a series of transactions are executed which update other tables.

+ Q1-1

2

- Write a program to perform the updates, using SQL **INSERT INTO** and **UPDATE** commands.
- Write the program as a single transaction bounded by **BEGIN TRANSACTION** and **COMMIT/END TRANSACTION** statements.

+ Transaction

3

■ Problems for update queries

- What happens when two queries update the same data item?
- How to deal with system failure?

■ Transaction

- A sequence of read and write operations
- With computation steps
- Termination
 - Commit
 - Abort

```
Begin_transaction Reservation
begin
  input(flight_no, date, customer_name);
  EXEC SQL UPDATE FLIGHT
      SET      STSOLD = STSOLD + 1
      WHERE FNO = flight_no
      AND     DATE = date;
  EXEC SQL INSERT
      INTO     FC(FNO,DATE,CNAME,SPECIAL)
      VALUES (flight_no,date,customer_name, null);
  output("reservation completed")
end.
```

+ Transaction Property

4

- Atomicity

- All or Nothing

- Consistency

- Correctness, from one valid state to another

- Isolation

- Cannot reveal its results to others before commitment

- Durability

- After commitment, results are permanent





```
1 BEGIN TRANSACTION
2 read inputfile EventIdent // EventIdent(Compldent, Pos)
3 while inputfile not empty
4     read inputfile Compldent, Pos // Tennis(Murray, 1)
5     INSERT INTO Results(EventID, ComplID, Position)
6     VALUES (EventIdent, Compldent, Pos)
7
8     if Pos < 4 then // Results(Tennis Man Single, Murray, 1)
9         UPDATE Medalists(ComplID, NMedals) // Medalists(Murray, 1)
10        SET NMedals = NMedals + 1
11        WHERE ComplID = Compldent
12
13        if Pos = 1 then
14            MedalAwarded = "Gold"
15        else if Pos = 2 then
16            MedalAwarded = "Silver"
17        else
18            MedalAwarded = "Bronze" // Medals (Tennis,Gold, Murray)
19
20        INSERT INTO Medals (EventID, Medal, ComplID)
21        VALUES (EventIdent, MedalAwarded, Compldent)
22
23        UPDATE MedalTally(Country, Medal, Number)
24        SET Number = Number + 1
25        WHERE MedalTally.Country =
26            (SELECT Country FROM Competitors
27             WHERE ComplID = Compldent)
28            AND Medal = MedalAwarded
29
30 COMMIT
31 END TRANSACTION // MedalTally(GBR, Gold, 12)
```

Tennis Man's single	
Andy Murray	1
Juan Martin Del Potro	2
KEI NISHIKORI	3
RAFAEL NADAL	4

+ Q 1-2 Issues

6

1. **INSERT INTO Results**(EventID, CompID, Position) → Add new ones
2. **UPDATE Medalists**(CompID, NMedals) → Update existing ones
3. **INSERT INTO Medals** (EventID, Medal, CompID) → Update existing ones
4. **UPDATE MedalTally**(Country, Medal, Number) → Update existing ones

Transaction 1: 200m butterfly(Michael Phelps, 1)

Transaction 2: 4*100m freestyle(Michael Phelps, 1)



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

7

So we need locks!

+ Q 1-3 Lock

8

- Q: Annotate your program with *read-lock*, *write-lock* and *unlock* statements.
 1. The notation should make clear the level of granularity of the locking.
 2. In each case indicates whether the granularity of locking is more than strictly necessary.
 3. Justify your decision based on the characteristics of the application.



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q1-3 Lock

9

BEGIN TRANSACTION

```
WRITE LOCK TABLE Results //not strictly necessary
WRITE LOCK TABLE Medals //not strictly necessary
WRITE LOCK TABLE MedalTally //not necessary at table level
WRITE LOCK TABLE Medalists //not necessary at table level
READ LOCK TABLE Competitors //not strictly necessary
```

// perform the body of the transaction in (a)

```
UNLOCK TABLE Results
UNLOCK TABLE Medals
UNLOCK TABLE MedalTally
UNLOCK TABLE Medalists
UNLOCK TABLE Competitors
```

COMMIT

END TRANSACTION



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA


+ Q1-4 2PL

10

■ **Q:** Show that your lock/unlock annotations constitute two-phase locking.

- No transaction should request a lock after it releases one of its locks
 - A transaction should not release a lock until it is certain that it will not request another lock
- *Growing Phase*: obtain locks and access data
- *Shrinking Phase*: release locks

■ **A:** ~



+ Q1-4 Dead Lock

11

1. Deadlock occurs on the tables need lock { **MedalTally**(Country, Medal, Number)
Medalists(CompID, NMedals)
2. Transactions do not obtain all the locks they need

T1:

WRITE LOCK TABLE **MedalTally**

//waiting for Medalists available
UPDATE Medalists(CompID, NMedals)
SET NMedals = NMedals + 1
WHERE CompID = Compldent

SET Number = Number + 1
WHERE **MedalTally**.Country =
(SELECT Country FROM Competitors
WHERE CompID = Compldent)
AND Medal = MedalAwarded

UNLOCK TABLE **MedalTally**

T2:

WRITE LOCK TABLE **Medalists**

UPDATE Medalists(CompID, NMedals)
SET NMedals = NMedals + 1
WHERE CompID = Compldent

//waiting for MedalTally available

SET Number = Number + 1
WHERE **MedalTally**.Country =
(SELECT Country FROM Competitors
WHERE CompID = Compldent)
AND Medal = MedalAwarded

UNLOCK TABLE **Medalists**

Please note: 2PL can make sure our results are correct.

This is a sufficient condition

But 2PL can trigger Dead Lock as well.

Example: Lecture Notes @Week 4 Page 18: T₁' & T₂'

+ Q1-5 Failures

12

- **Q:** Is there any circumstance in which the program would have to abort the transaction?
- **Transaction Failures**
 - Incorrect input data
 - Present/Potential deadlock
 - Data are accessed by another transaction
- **System Failures**
 - Media failure, processor failure, communication break, power outage...
- **Our Program?**
 - Why we need DBMS?
 - Error handling



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2 Distributed Replication

13

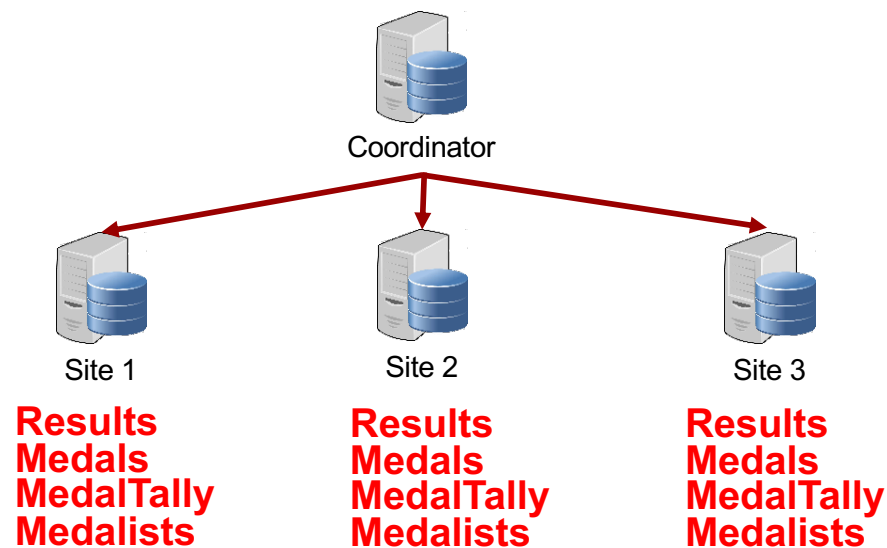
- Each table in the above question can be stored at a different site, and some tables, such as **Medals**, **Medalists** and **MedalTally** may need to be replicated at several sites.
- Clearly, the transaction in the previous question will comprise of a set of **sub-transactions** executing at different sites to updating data there.
- The transaction can only finish when all its sub-transactions successfully finish.
- This process is enforced by using the two-phase commit (2PC) protocol.



+ Q2 Distributed Replication

14

- For each Transaction
 - Break into sub-transactions for each site



+ Q2 Distributed Data Replication

15

■ Purposes

- System Availability
- Performance
- Scalability
- Application Requirements

■ Replicas converge to the same value

■ Issue

- Consistency
 - Where updates are performed
 - How the updates propagate



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-1 Synchronous Replication

16

- When the update commits, all the copies have the same value
 - Benefits
 - Consistent
 - No need for remote read
 - Easy to handle failure
 - Drawbacks
 - A transaction has to update all the copies before it can terminate
 - The **response time** performance suffers from the slowest one
 - If **any is unavailable**, then the whole transaction cannot commit



+ Q2-1 Synchronous Replication

17

- **Q:** Suppose there are a large number of replicas for **MedalTally**. Should we adopt the synchronous replication strategy for this table?



- All copies of the table should be **consistent** with each other at all time.
- The more sites, the higher possibility of failure \ Locks
- If one of the replica sites fails during the execution of a transaction with update operations, the whole transaction needs to be aborted.

+ Q2-2 Asynchronous Replication

18

- Partially updated
 - Tolerate **inconsistency** for better **performance**
 - Lower response time for update transaction
 - Not all consistent, some replicas may be **out-of-date**
 - Eventually consistency




THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-2 Asynchronous Replication

19

- Q: Would you recommend using asynchronous replication for this application?
- Read
 - Tolerable
- Update
 - Need strategies



+ Q2-3 Distributed Reliability

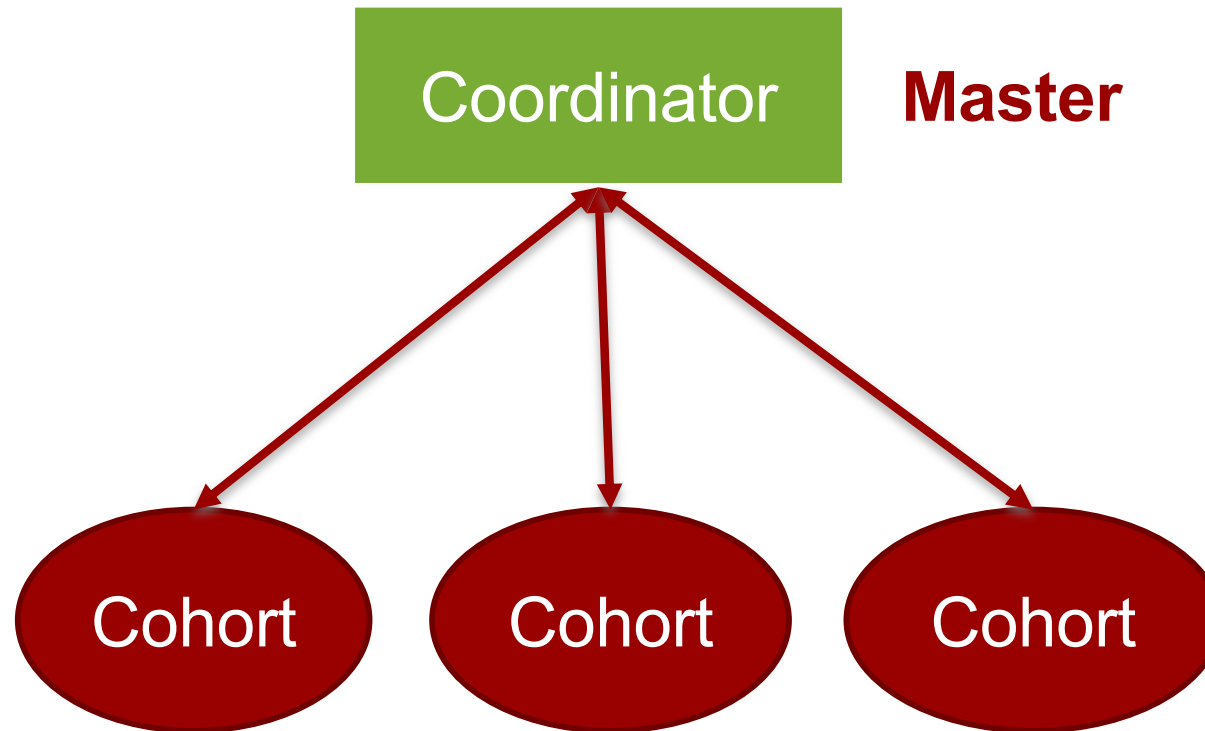
20

- **Q:** Suppose one of the **MedalTally** replica sites fails during the transaction. Show the exchange of messages among sub-transactions resulting in the two-phase commit (2PC) protocol to issue an abort instruction.
 - Atomicity and Durability
 - All or Nothing
 - Commit then permanent
 - All sites involved in the execution of a distributed transaction **agree to commit** before its effects are made **permanent**



+ Q2-3 Distributed Database Structure

21



+ Q2-3 2PC

22

- Protocol

- TCP/IP, HTTP, SSH ...
- Two Phase Commit Protocol

1. **Commit-Request Phase** (or Voting Phase)

1. Coordinator ask cohorts to prepare/work
2. Cohorts response yes/no
 - Does not know the situation of the whole transaction, only about itself
 - If responses yes, the cohort has to wait for the response from coordinator.
 - Cannot abort by itself

2. **Commit Phase**

3. If all yes, ask cohorts to commit
 - Else, tell cohorts to abort
4. Cohorts ACK to coordinator



+ Q2-3 Success

23

//execution/Preparation phase

Coordinator prepare to UpdateResults
Coordinator prepare to UpdateMedalists
Coordinator prepare to UpdateMedals
Coordinator prepare to UpdateMedalTally

// prepare-to-commit phase

UpdateResults yes to Coordinator
UpdateMedalists yes to Coordinator
UpdateMedals yes to Coordinator
UpdateMedalTally yes to Coordinator

//commit-or-abort phase

Coordinator commit to UpdateResults
Coordinator commit to UpdateMedalists
Coordinator commit to UpdateMedals
Coordinator commit to UpdateMedalTally

//execution status reporting phase

UpdateResults ack to Coordinator
UpdateMedalists ack to Coordinator
UpdateMedals ack to Coordinator
UpdateMedalTally ack to Coordinator



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-3 Failure

24

//execution/Preparation phase

Coordinator prepare to UpdateResults
Coordinator prepare to UpdateMedalists
Coordinator prepare to UpdateMedals
Coordinator prepare to UpdateMedalTally

// prepare-to-commit phase

UpdateResults yes to Coordinator
UpdateMedalists yes to Coordinator
UpdateMedals yes to Coordinator
UpdateMedalTally no to Coordinator

//failed subtransaction

//commit-or-abort phase

Coordinator abort to UpdateResults
Coordinator abort to UpdateMedalists
Coordinator abort to UpdateMedals
Coordinator abort to UpdateMedalTally

//execution status reporting phase

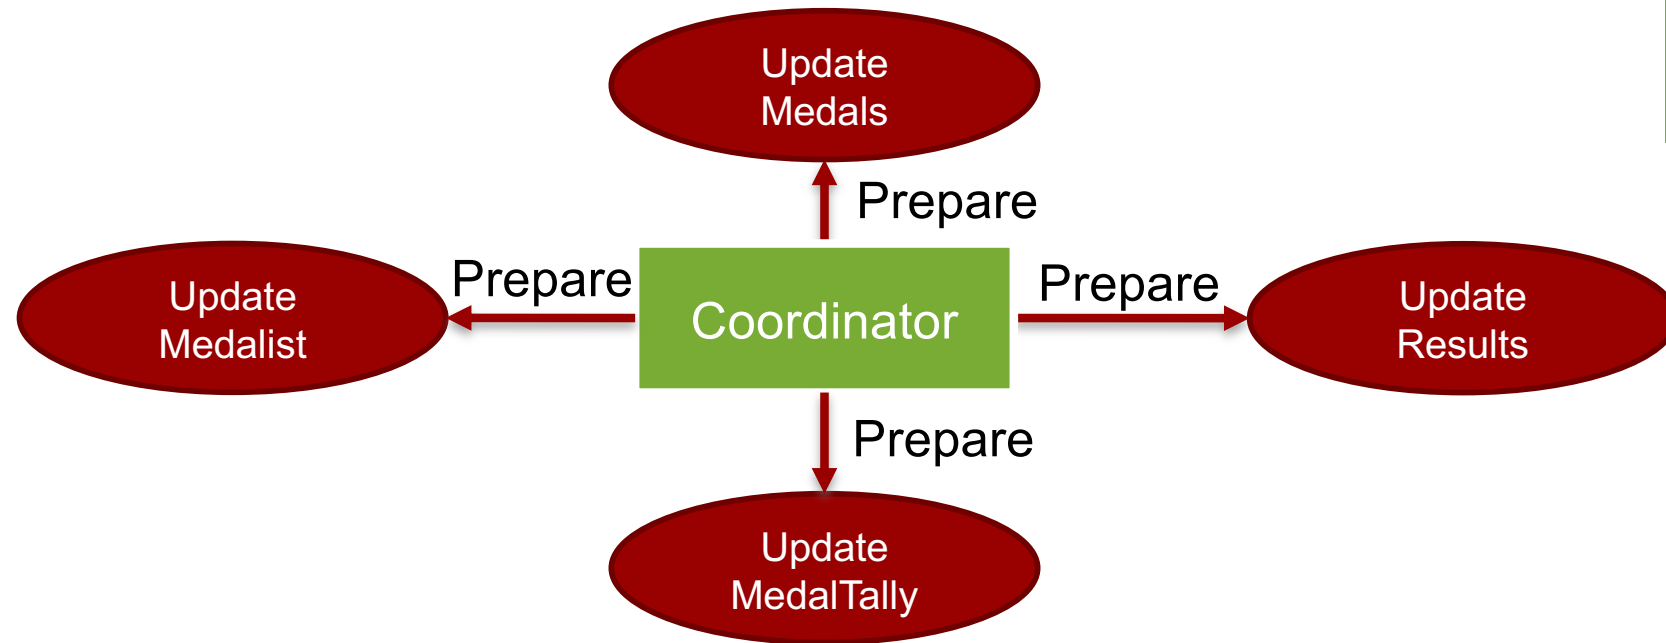
UpdateResults ack to Coordinator
UpdateMedalists ack to Coordinator
UpdateMedals ack to Coordinator
UpdateMedalTally ack to Coordinator



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-3 2PC Voting Phase 1

25



//Execution/Preparation phase

Coordinator execute to UpdateResults

Coordinator execute to UpdateMedalists

Coordinator execute to UpdateMedals

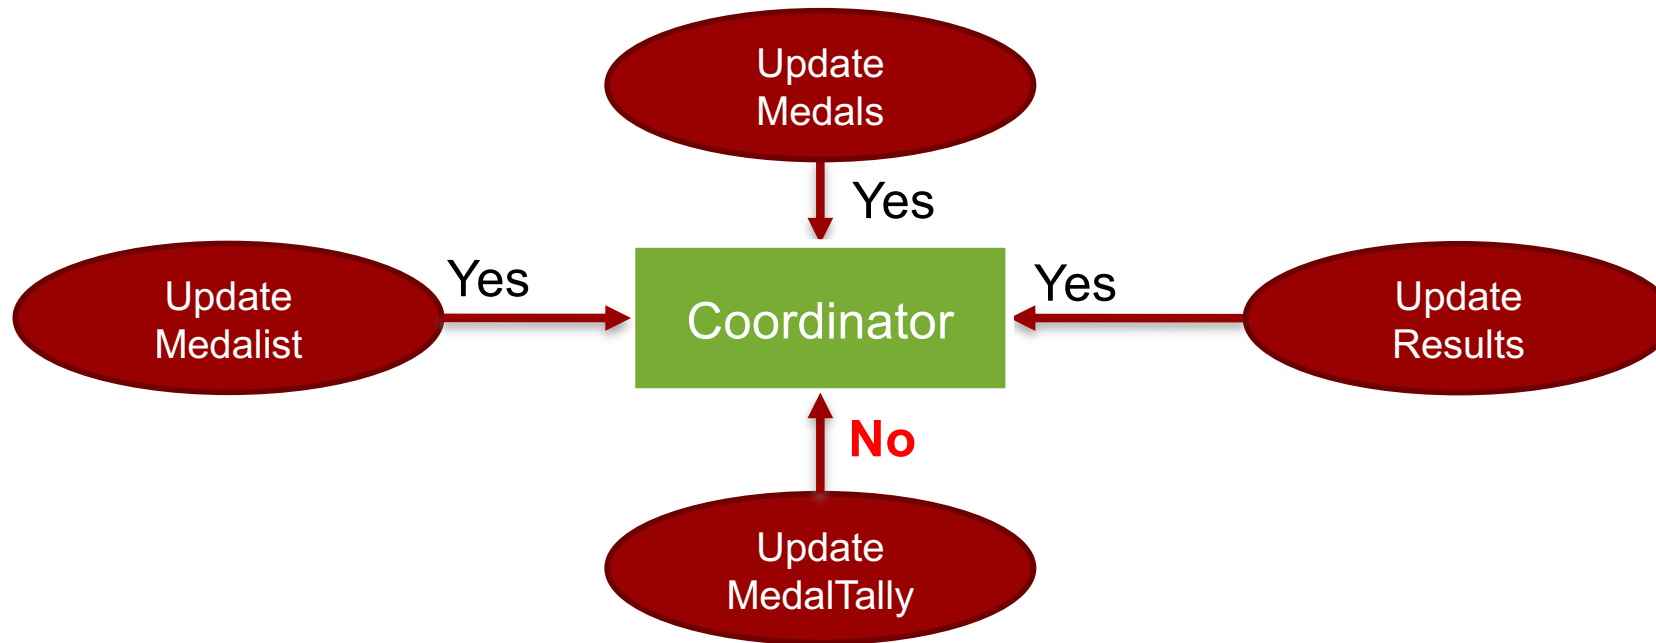
Coordinator execute to UpdateMedalTally



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-3 2PC Voting Phase 2

26



// prepare-to-commit phase

UpdateResults **yes** to Coordinator

UpdateMedalists **yes** to Coordinator

UpdateMedals **yes** to Coordinator

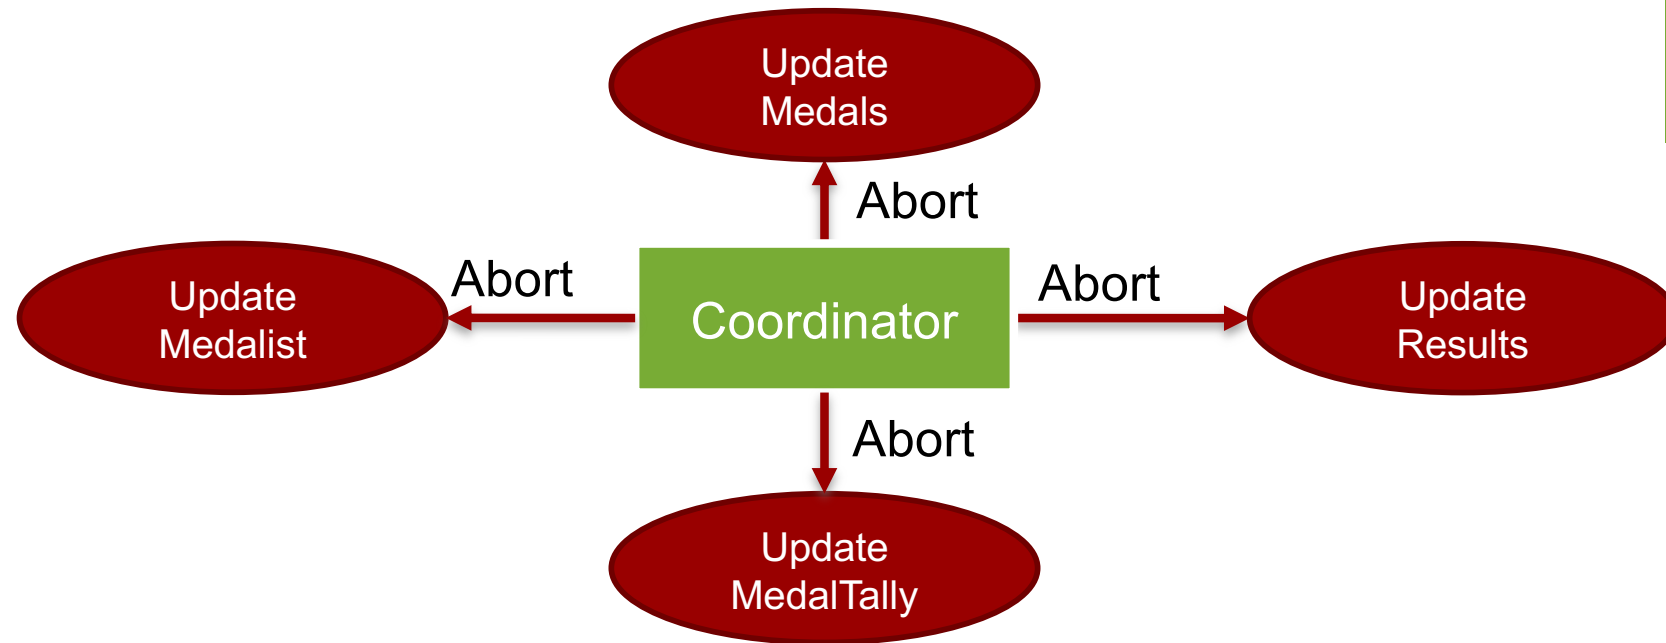
UpdateMedalTally **no** to Coordinator

//failed subtransaction

- All or Nothing
 - If all yes, commit
 - If any no, abort

+ Q2-3 2PC Commit/Abort Phase 1

27



// prepare-to-commit phase

UpdateResults yes to Coordinator

UpdateMedalists yes to Coordinator

UpdateMedals yes to Coordinator

UpdateMedalTally no to Coordinator

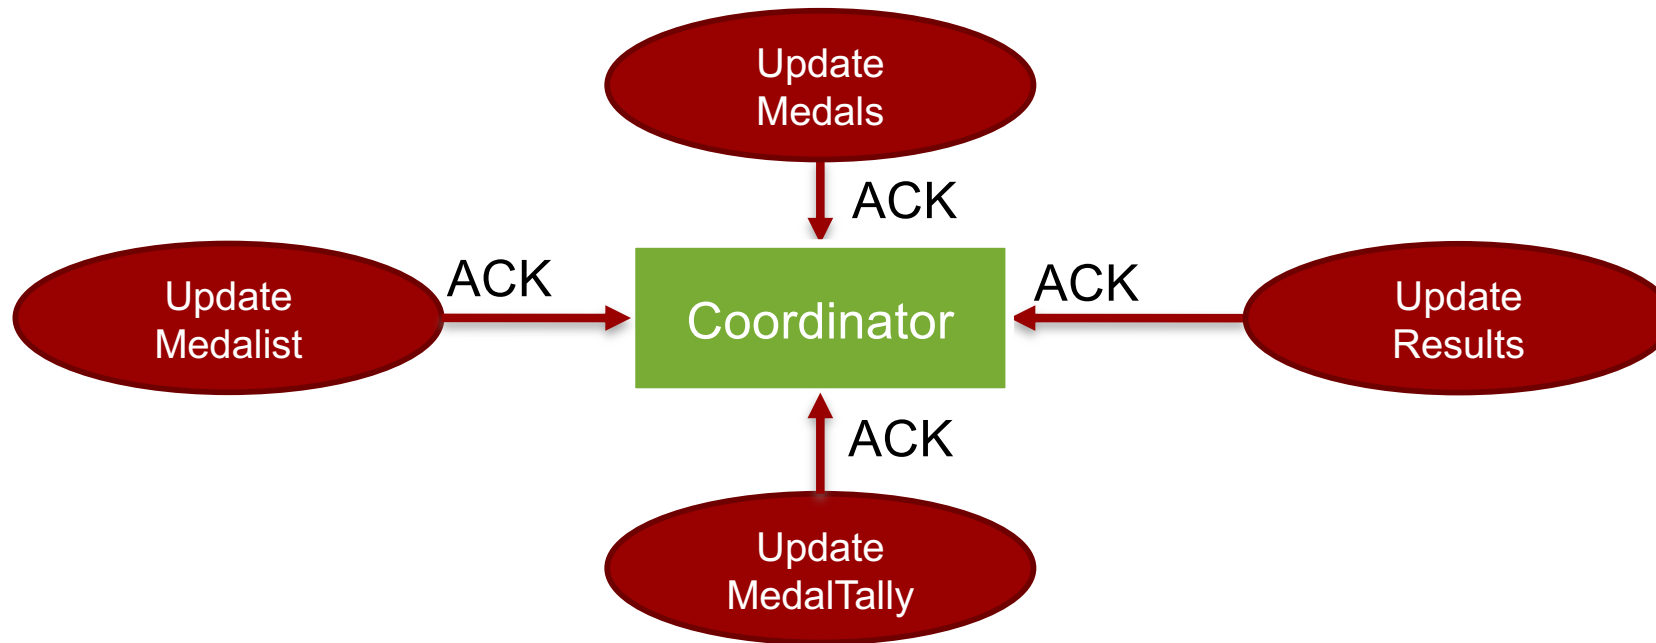
//failed subtransaction



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-3 Commit/Abort Phase 2

28



//execution status reporting phase

UpdateResults ack to Coordinator

UpdateMedalists ack to Coordinator

UpdateMedals ack to Coordinator

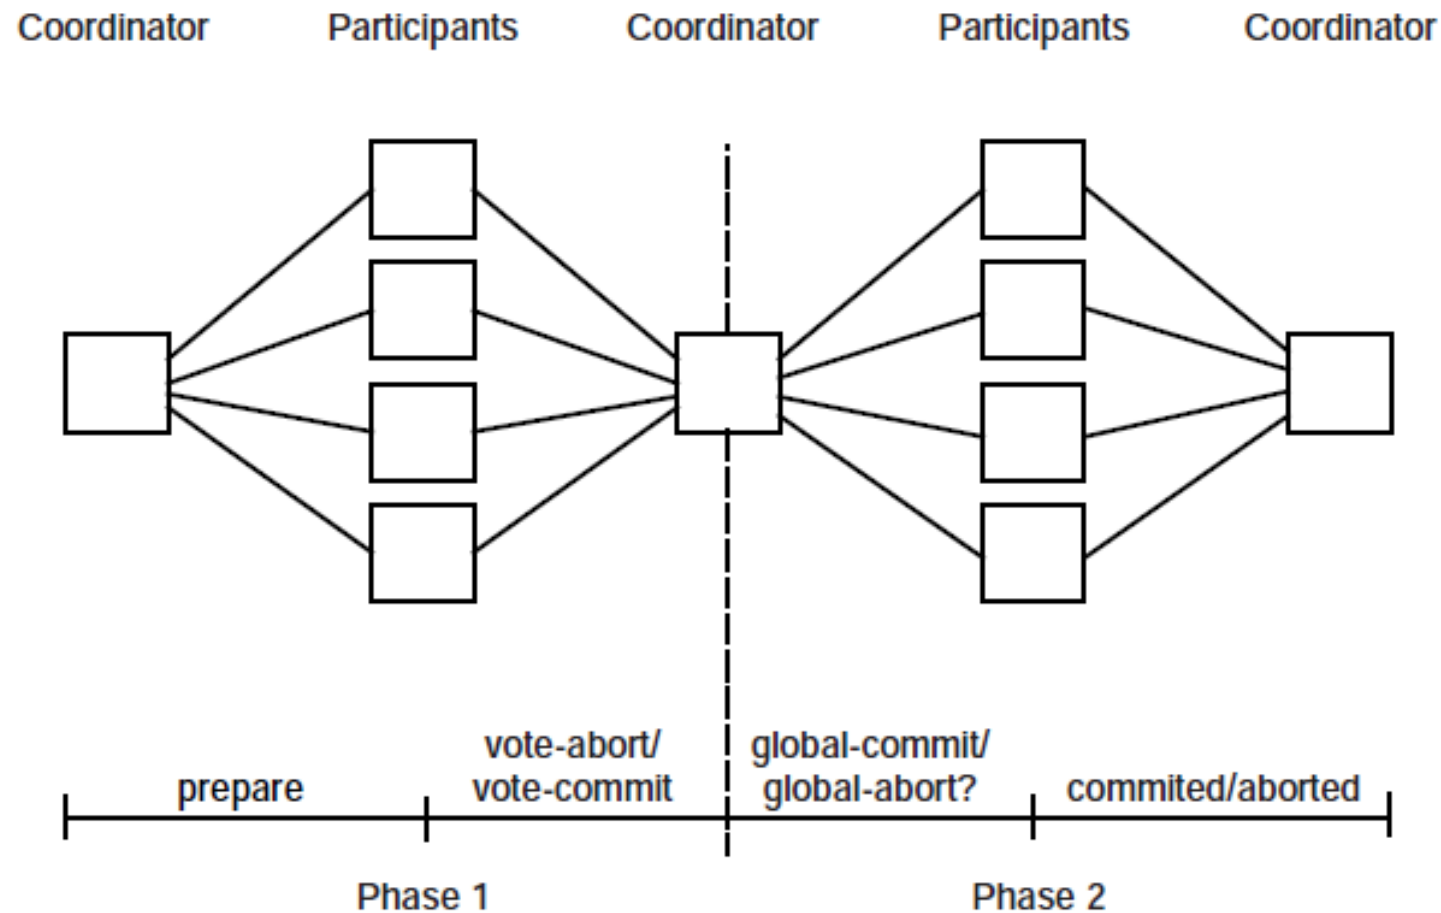
UpdateMedalTally ack to Coordinator



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q2-3 2PC Communication Structure

29



+ Q3 2PC Failure Handling

30

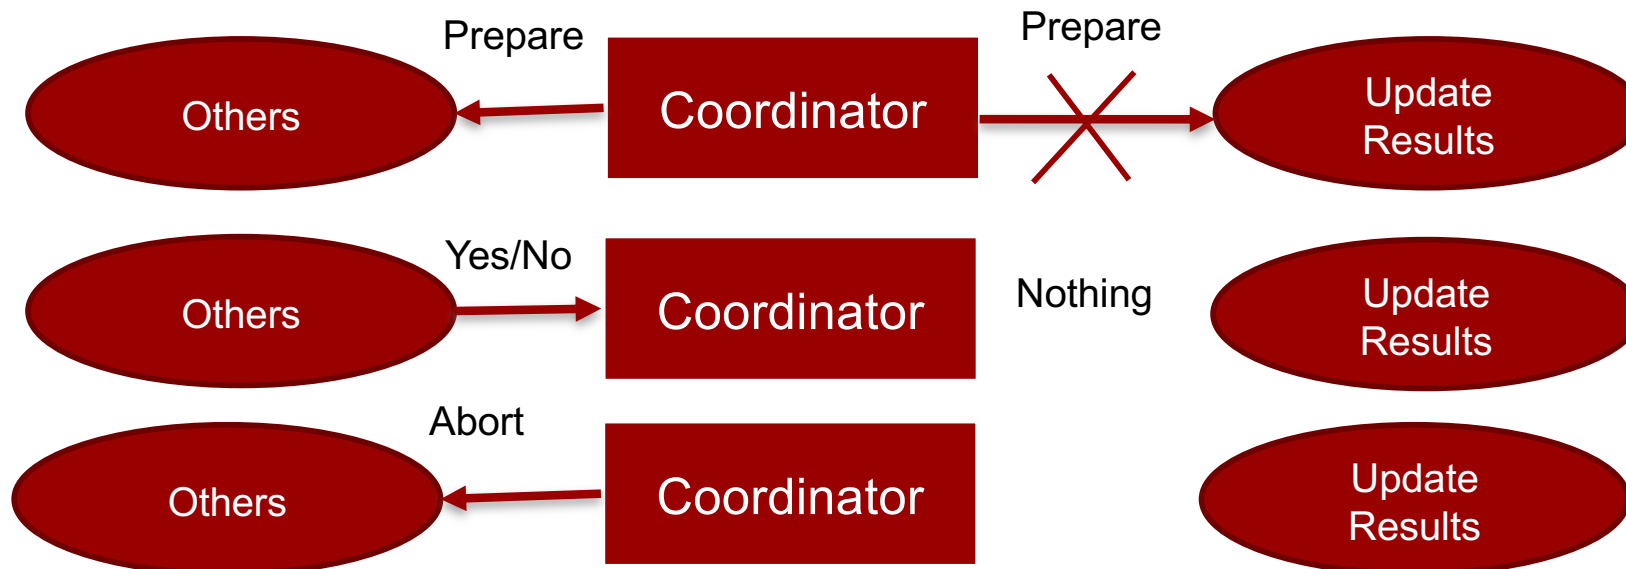
- When 2PC is used as the commit protocol, explain how the system recovers from failure and deals with particular *transaction T* in each of the following cases:
 - a. A **subordinate** site for *T* fails **before receiving** a *prepare* message.
 - b. A **subordinate** site for *T* fails **after receiving** a *prepare* message but **before making a decision**.
 - c. The **coordinator** site for *T* fails **before sending** a *prepare* message.
 - d. The **coordinator** site for *T* fails **after writing** an *abort* log record but **before sending** any further messages to its subordinates.



+ Q3-1 Cohort Fail

31

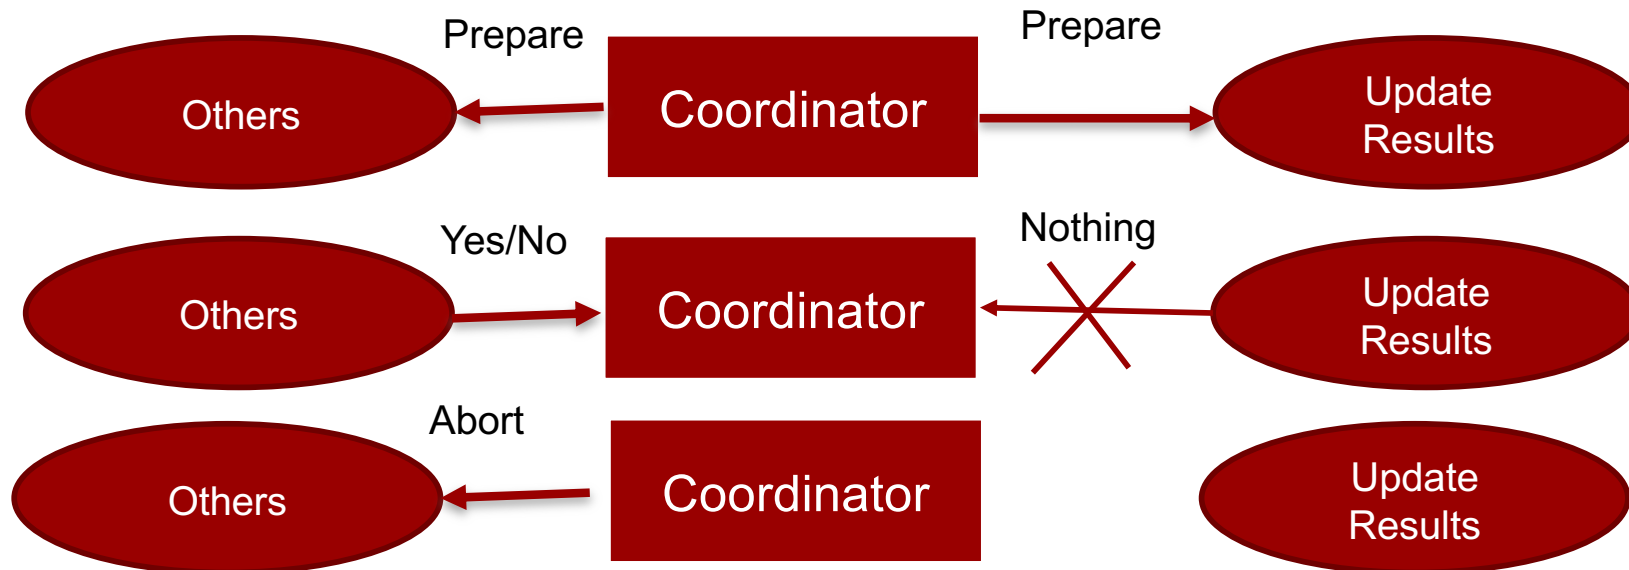
- Q: A subordinate site for T fails before receiving a prepare message
- A: The coordinator will not receive either a yes or a no message from the subordinate, and will abort the transaction (once time-outed).
- It will instruct all the subordinates that replied yes to abort.
- Once the subordinate recovers, it will check with the coordinator, which will send the default abort message and the sub transaction will also be aborted.



+ Q3-2 Cohort Fail

32

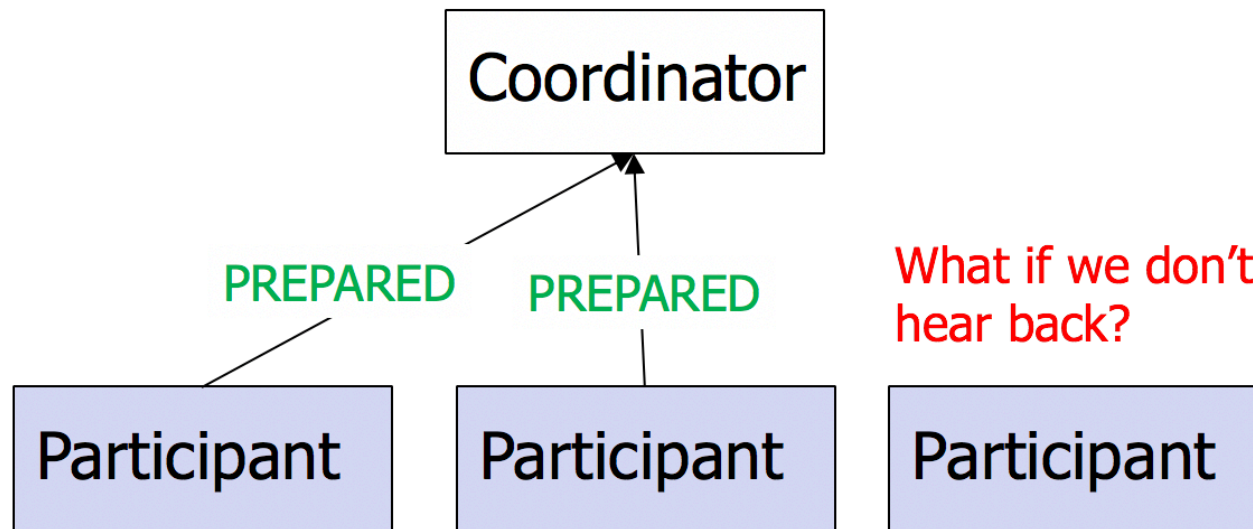
- **Q**: A subordinate site for **T** fails after receiving a *prepare* message but before making a decision
- **A**: The situation is the same as above, since the subordinate has not yet made a decision. No message has been sent to the coordinator.



+ Q3-2 Cohort Fail

33


- **Q** : A subordinate site for **T** fails after receiving a prepare message but before making a decision
- **A** : The situation is the same as above, since the subordinate has not yet made a decision. No message has been sent to the coordinator.



+ Q3-3 Coordinator Fail

34

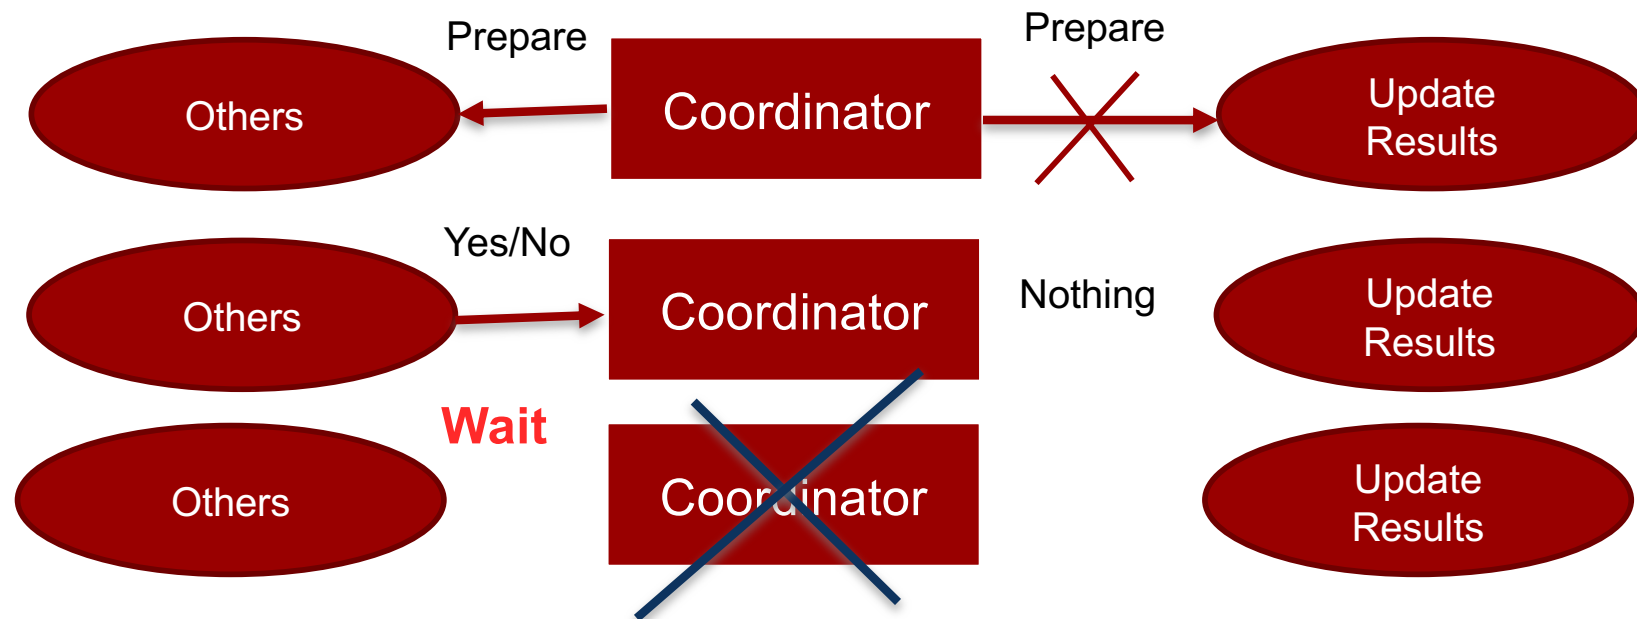
- Q: The **coordinator** site for **T** fails **before sending** a **prepare** message.
- A: Since there is no prepare, commit or abort log record, T can be unilaterally aborted and undone, and an end log record written. Since the coordinator did not send the prepare message before it crashed, the subordinates may abort once time-outed.



+ Q3-4 Coordinator Fail

35

- Q: The coordinator site for T fails after writing an *abort* log record but before sending any further messages to its subordinates.



Subordinates cannot abort unilaterally after voted yes.

So 2PC is a blocking protocol → 3PC

+ Q3 3PC

36

■ Assumptions

1. Each site uses the write-ahead-log protocol
2. At most one site can fail during the execution of the transaction
3. Set a timer for abort

■ 2PC

■ Phase 1:

- Cohorts know the transaction ahead
- Does not know others' conditions

■ Phase 2:

- Waiting for the next instruction from coordinator
- Commit/Abort

■ 3PC

■ Phase 1:

- Cohorts know the transaction ahead
- Does not know others' conditions

■ Phase 2:

■ PreCommit

- Cohorts know the transaction will start
 - Even if coordinator fails, no "Commit" received, it is OK to commit when time out
 - If one of the cohorts fails after voting "yes", coordinator does not have enough "ack", coordinator sends abort
 - If one of the cohorts fails after receiving "PreCommit", it still commits after reading its log
- Abort/Received nothing
 - Cohorts know the transaction will abort

■ Phase 3:

- Commit/Abort

+ Q3 3PC

37

Further reading: https://en.wikipedia.org/wiki/Three-phase_commit_protocol

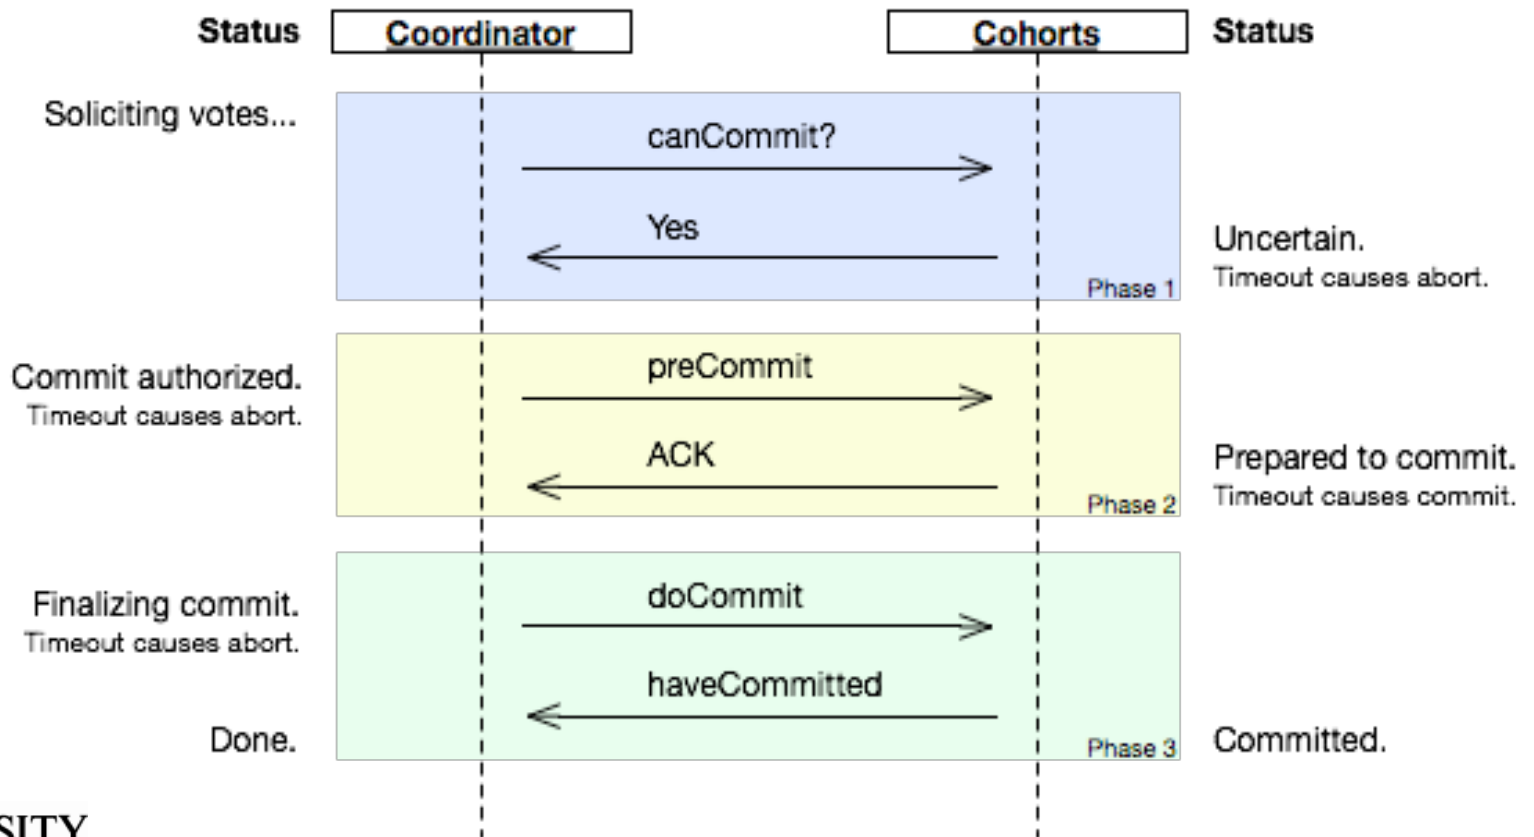
■ Non-blocking

- Places an upper bound on the amount of time required before a transaction either commits or aborts

Phase 1 in 2PC

New in 3PC

Phase 2 in 2PC



+ Q4 Distributed Data Consistency

38

- Distributed transaction management can use **primary-copy based** approach or **voting-based** approach to maintain **data consistency** among multiple copies.
- Please explain how these two approaches work, and compare their advantages and disadvantages.




THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA

+ Q4 Primary Copy

39

- The **primary-copy approach** selects one copy as the *master copy* (e.g., the one at the 'birth site' of the data item).



- Simple to update
- Primary Site
 - Primary site overload
 - Primary Copy
- Others Temporary Inconsistency

+ Q4 Voting Based

40

■ Version Number

- Monotonically increasing

■ N copies

- Write at least m
- Read at least n
- $m + n > N$
- Guarantee can read the latest version

Version No.

....

0102

0103

0104

$m+n > \text{Total number}$

Write: 3
(m)


Read: ?
(n)

■ Advantages

- More **reliable**
- Writes fewer copies.

■ Disadvantages

- Read transaction needs to **check many sites**



THE UNIVERSITY
OF QUEENSLAND
AUSTRALIA