INFS7901
Database Principles

Binary Search Trees

Hassan Khosravi & Ash Rahimi

Notes

- Notebooks so far:
  - Asymptotic analysis
  - Sorting
  - Trees

- Project:
  - Read the project description and deadline for final submissions.

- RiPPLE:
  - Third round deadline was today, Check fourth round due

- Final exam (read the ECP)
  - Only two hours, Blackboard test, using ProctorU.

Learning Outcomes

Description	Tag
Describe the different operations that are supported by the dictionary ADT.	Dict-ADT
Demonstrate how arrays can be used for implementation of a dictionary ADT.	Dict-Array
Demonstrate how linked lists can be used for implementation of a dictionary ADT.	Dict-Linkedlist
Describe the properties of binary trees.	Dict-Trees
Describe the properties of Binary Search Trees (BSTs).	
Represent BSTs using linked data structures.	
Insert elements into a BST.	
Delete elements from a BST.	
search for an element in a BST.	
Determine if a given tree is an instance of a BST.	
Compare and contrast algorithms for storing and subsequently retrieving values from a dictionary ADT.	
Reflect on how BSTs and their extensions can be used in DBMS.	

Dictionary ADT
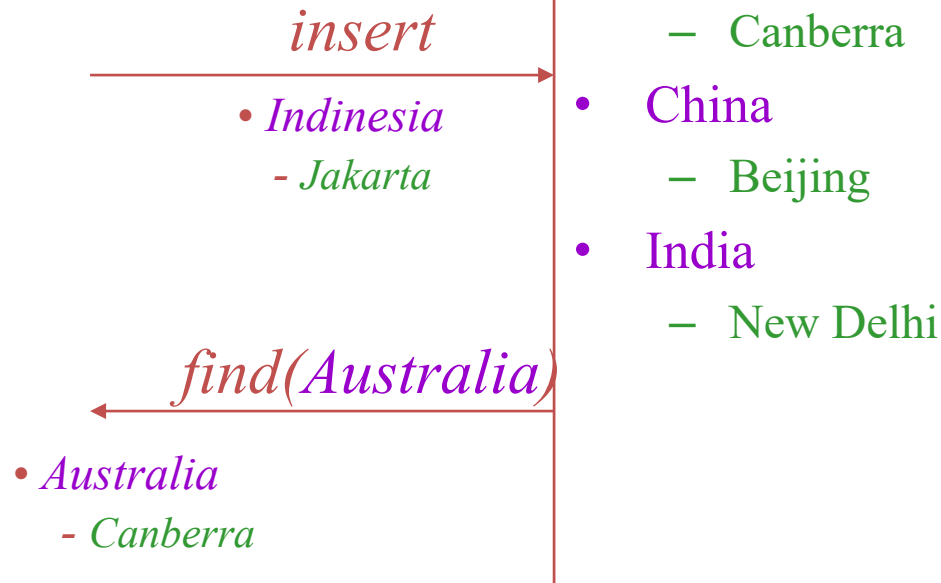
Tree Terminology

Binary Search Trees (BSTs)

Insertion and Deletion in BSTs

Dictionary Abstract Data type (ADT)

- Dictionary operations
  - create
  - destroy
  - insert
  - find
  - Delete

insert

→

- *Indinesia*
  - *Jakarta*

find(Australia)

←

- *Australia*
  - *Canberra*

- Australia
  - Canberra
- China
  - Beijing
- India
  - New Delhi

- Stores *values* associated with user-specified *keys*
  - values may be any (homogenous) type
  - keys may be any (homogenous) comparable type

Search/Set ADT

- Dictionary operations
  - create
  - destroy
  - insert
  - find
  - Delete

- Stores keys
  - keys may be any (homogenous) comparable
  - quickly tests for membership

- Australia
- China
- India
- Indonesia

insert
- *Japan*

find(Malaysia)
NOT FOUND

Naïve Implementations

- Unsorted Array

2	4	11	98	37	44	3			

Operation	Implementation	Running Time
Insert	Add after the current last value, if there is space available.	O(1)
Find	Scan the entire array	O(n)
Delete (index known)	Delete item, fill empty space with the last element	O(1)
Delete (value known)	Scan the entire array to find the item then delete	O(n)

Naïve Implementations

- Sorted Array

2	3	4	11	37	44	98			

Operation	Implementation	Running Time
Insert	Shift to make room for the insertion	O(n)
Find	Use binary search	O(log n)
Delete (index known)	Shift to get rid of empty space	O(n)
Delete (value known)	Use binary search + shift	O(n)

Linked Lists

- Consider the following <u>abstraction</u>, picturing a short linked list:

N	5			1			2	N
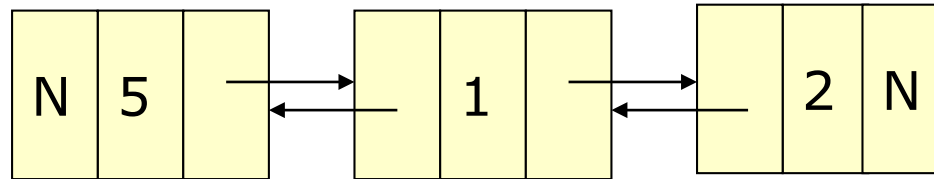
- What might it look like in memory?

	data	*previous*	*next*
1080	2	100	Null
…			
600	5	Null	100
…			
140			
120			
100	1	600	1080

Node
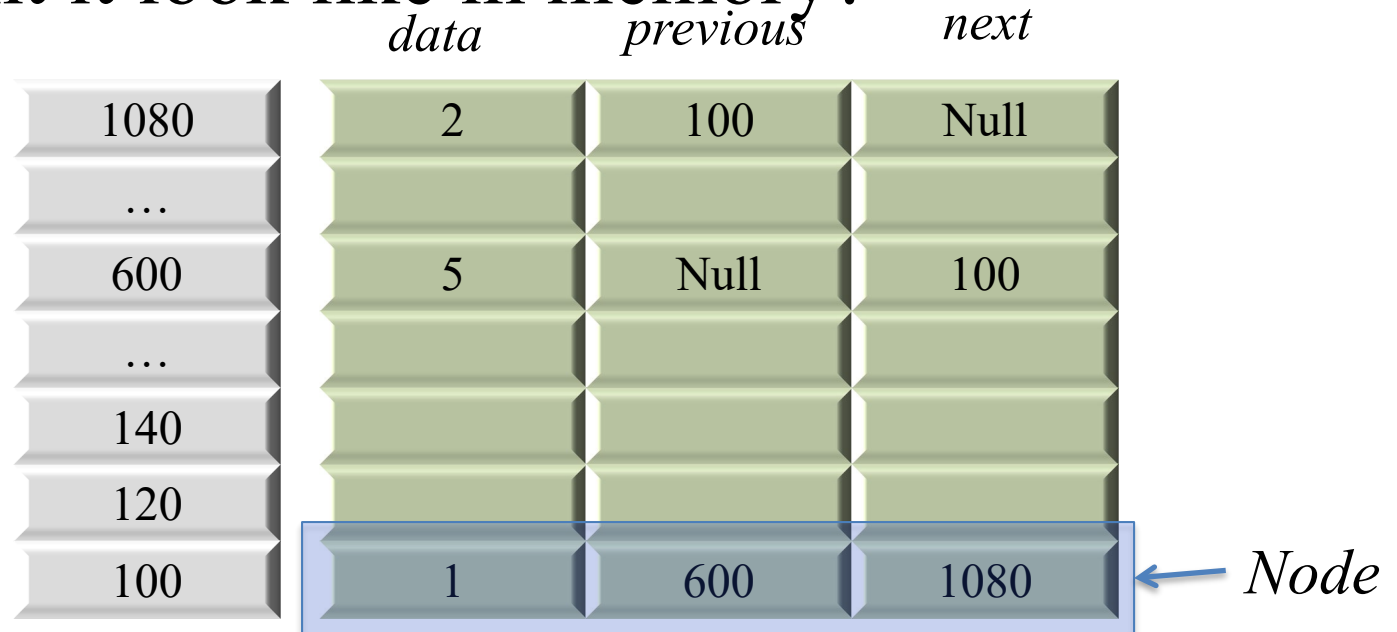
Inserting an Element to a Linked List
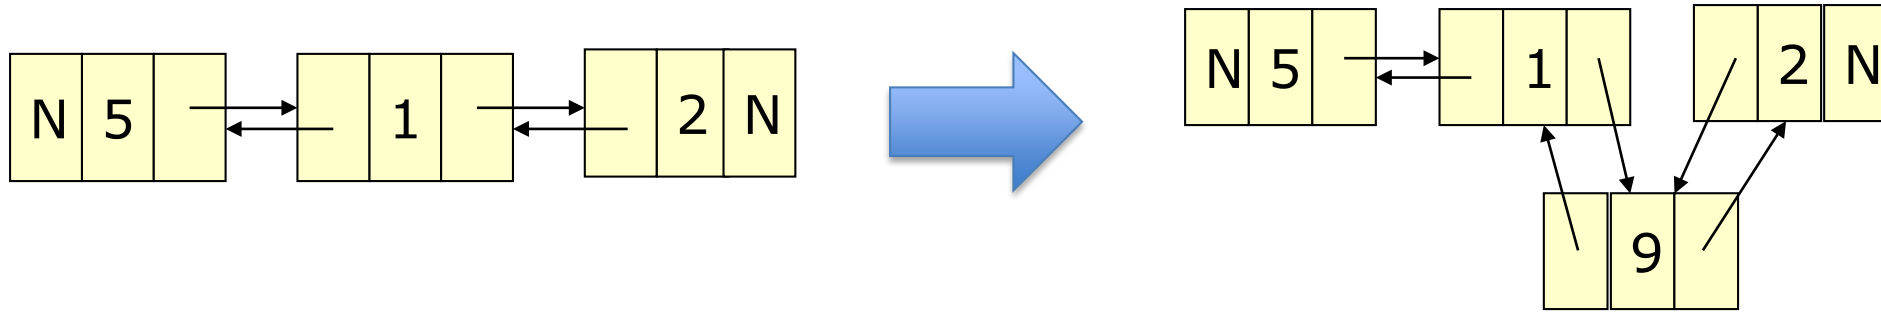
	data	previous	next
1080	2	140	Null
…			
600	5	Null	100
…			
140	9	100	1080
120			
100	1	600	140

Removing an Element from a Linked List
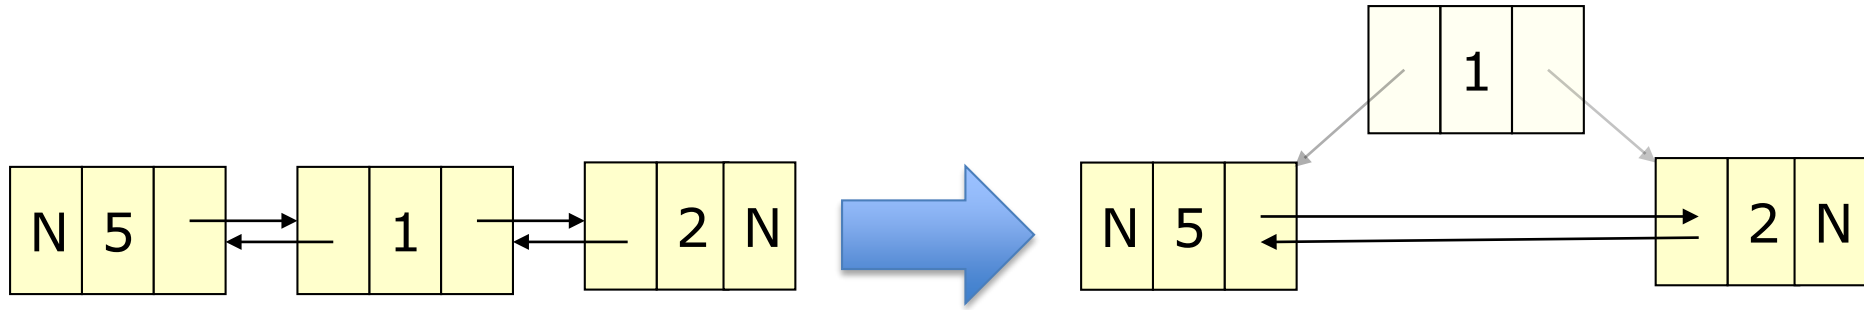
	data	previous	next
1080	2	600	Null
…			
600	5	Null	1080
…			
140			
120			
100	1	600	1080	← *delete*

https://visualgo.net/en/list

Naïve Implementations

- Unsorted linked list
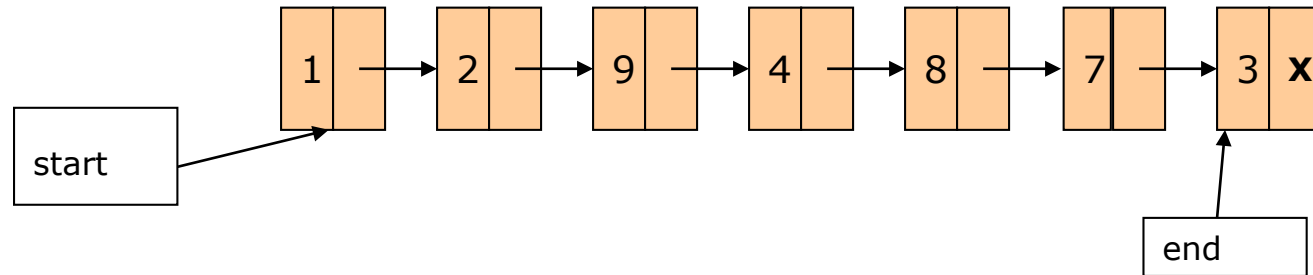
Operation	Implementation	Running Time
Insert	Add after end	O(1)
Find	Scan list	O(n)
Delete (Address known)	Remove and rechain	O(1)
Delete (value known)	Scan list + remove and rechain	O(n)

Naïve Implementations

- Sorted linked list
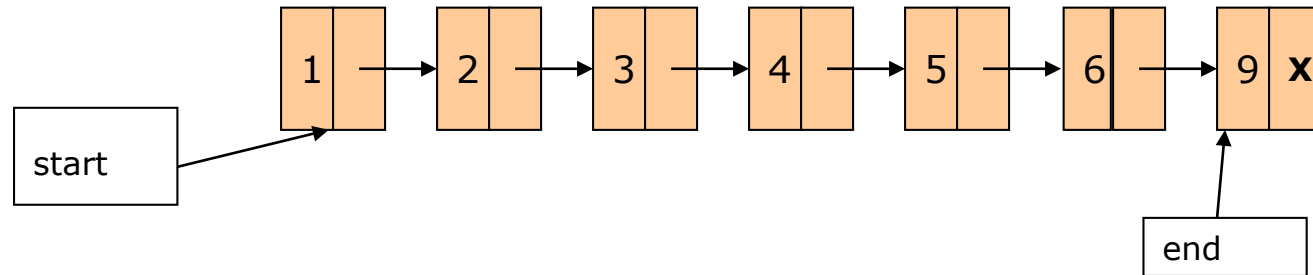
Operation	Implementation	Running Time
Insert	Scan list to find where to insert	O(n)
Find	Scan list	O(n)
Delete (Address known)	Remove and rechain	O(1)
Delete (value known)	Scan list + remove and rechain	O(n)

Naïve Implementations

	insert	*find*	*delete* By value	*delete* By address
Linked list				
– Unsorted	*O(1)*	*O(n)*	*O(n)*	*O(1)*
– Sorted	*O(n)*	*O(n)*	*O(n)*	*O(1)*
Array				
– Unsorted	*O(1)*	*O(n)*	*O(n)*	*O(1)*
– Sorted	*O(n)*	*O(lg n)*	*O(n)*	*O(n)*
Can we do better?	O(lg n)	O(lg n)	O(lg n)	O(lg n)

From Binary Search to Binary Search Trees
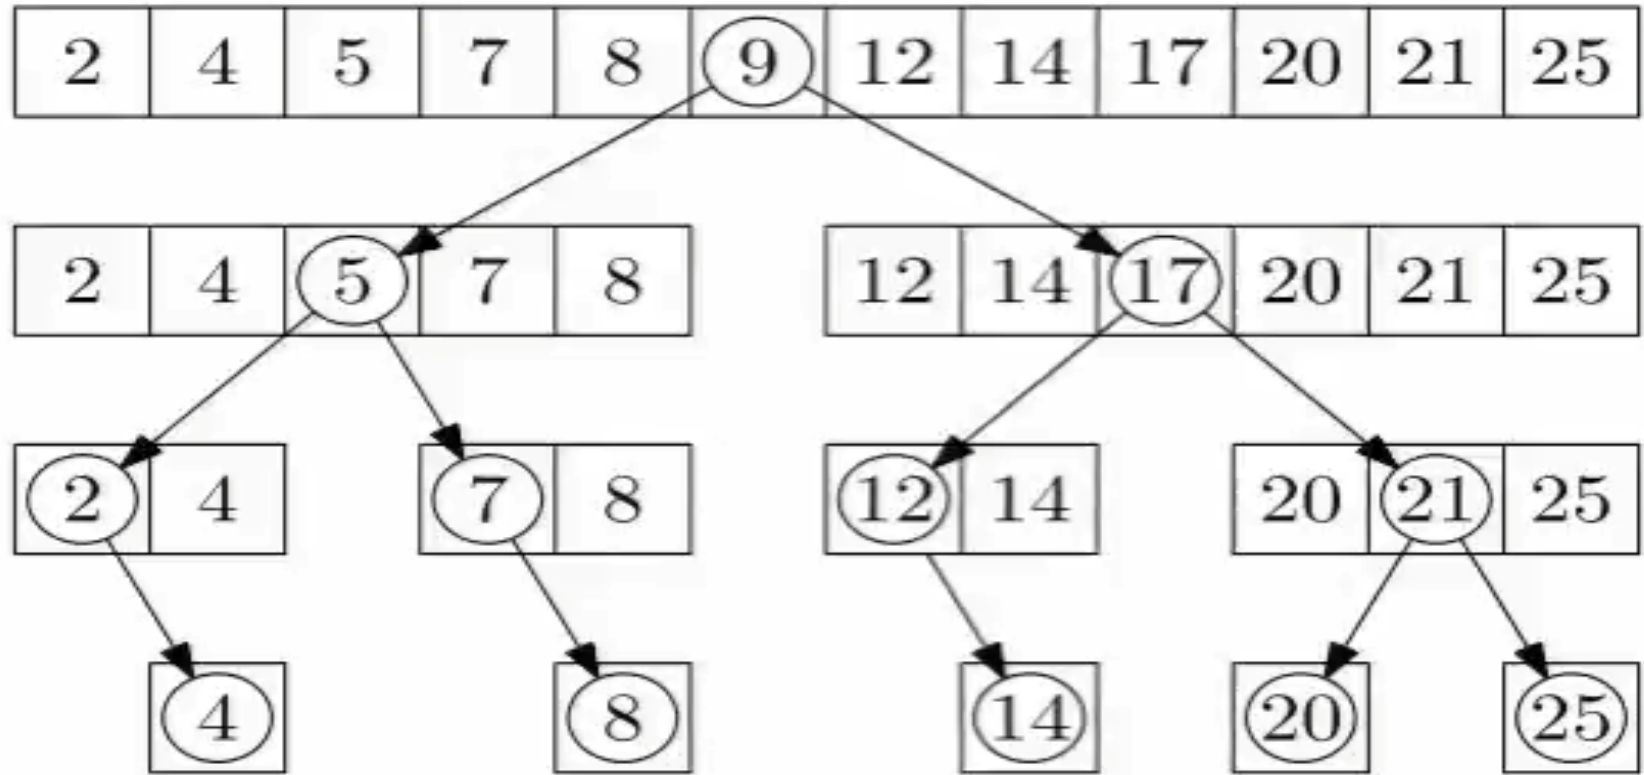
General idea (informally): You are allowed to shove new nodes into your list without having to shift things over.

Dictionary ADT

Tree Terminology

Binary Search Trees (BSTs)

Insertion and Deletion in BSTs

Tree Terminology

- *root:* the single node with no parent
- *leaf:* a node with no children
- *child:* a node pointed to by me
- *parent:* the node that points to me
- *Sibling:* another child of my parent
- *ancestor:* my parent or my parent's ancestor
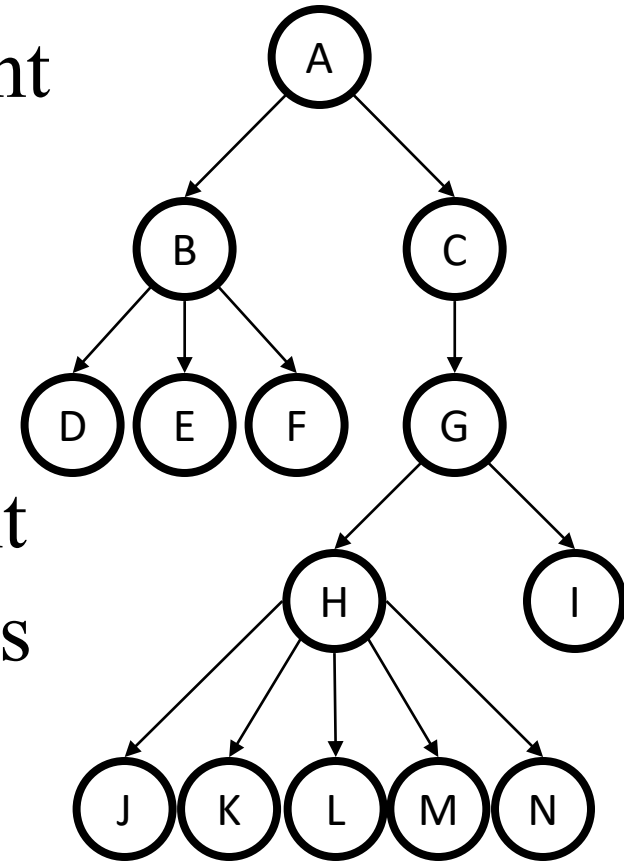
- *descendent:* my child or my child's descendent
- *subtree:* a node and its descendants

Tree Terminology

- *depth:* # of edges along path from root to node
  - *depth of H?*
    - *3*

- height: # of edges along longest path from node to leaf or, for whole tree, from root to leaf
  - height of tree?
    - 4
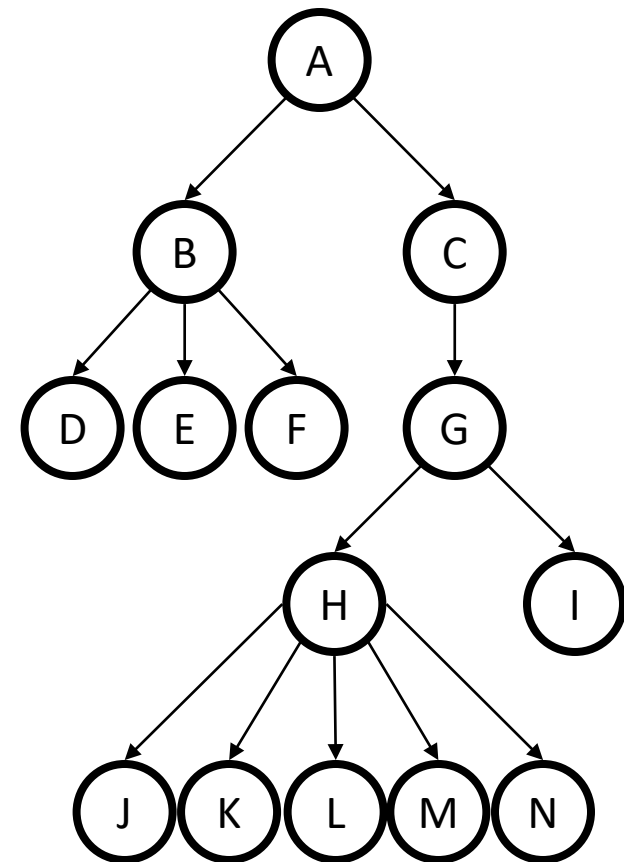
Tree Terminology

- *degree:* # of children of a node
  - *degree of B?*
    - *3*

- branching factor: maximum degree of any node in the tree

2 for binary trees,
5 for this weird tree
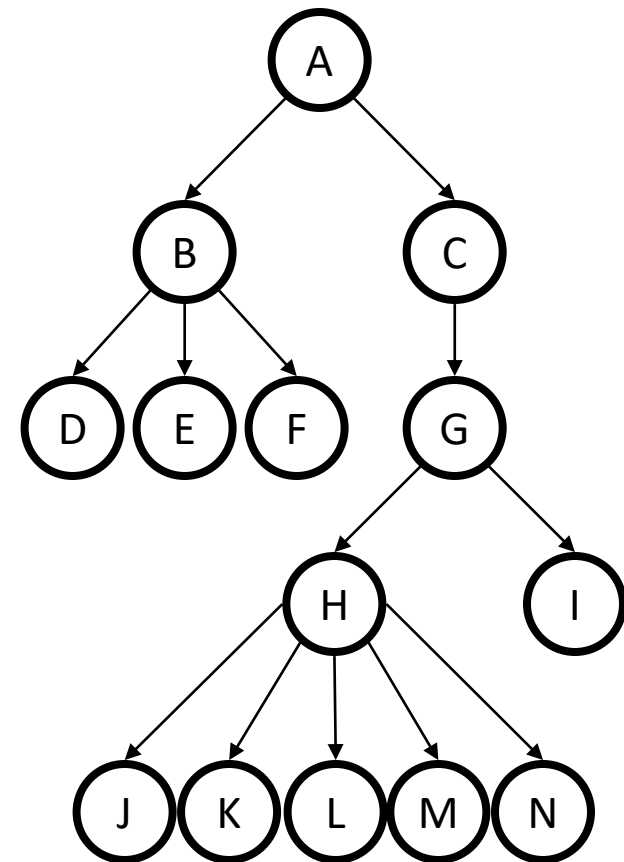
One More Tree Terminology Slide

- binary: branching factor of 2 (each child has at most 2 children)

- n-ary: branching factor of n

- complete: "packed" binary tree; as many nodes as possible for its height
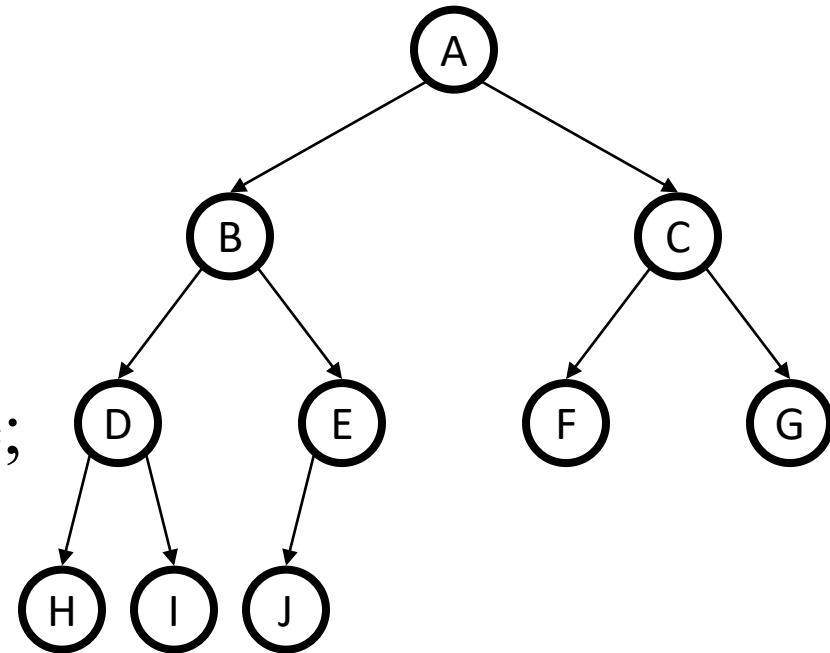
- nearly complete: complete plus some nodes on the left at the bottom

Trees and (Structural) Recursion

A tree is either:

- the empty tree
- a root node and an ordered list of subtrees

Trees are a recursively defined structure, so it makes sense to operate on them recursively.

Dictionary ADT

Tree Terminology

Binary Search Trees (BSTs)

Insertion and Deletion in BSTs

Binary Search Tree

- ## Binary tree property
  - each node has $\leq$ 2 children
  - result:
    - operations are simple

```
          8
         / \
        5   11
       / \  / \
      2  6 10  12
       \  \  \   \
        4  7  9   14
                   \
                    13
```
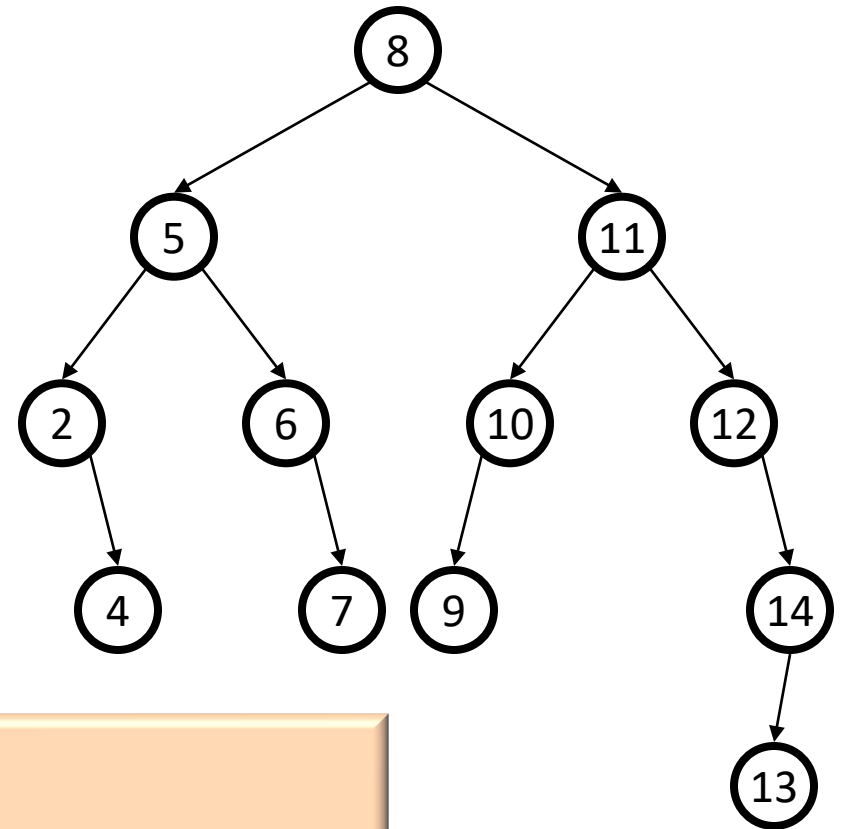
- ## Search tree property
  - **all** keys in left subtree smaller than root's key
  - **all** keys in right subtree larger than root's key
  - result:
    - easy to find any given key

Example and Counter-Example
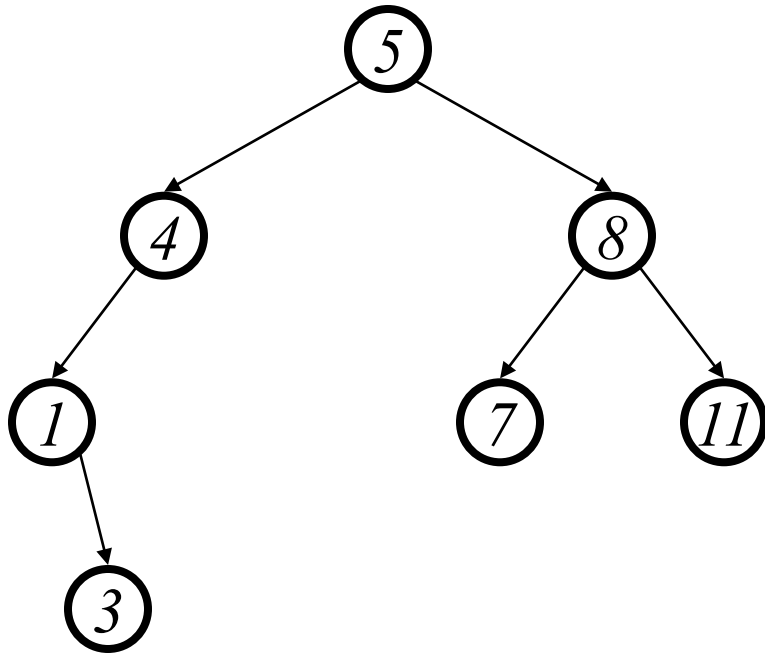
BINARY SEARCH TREE
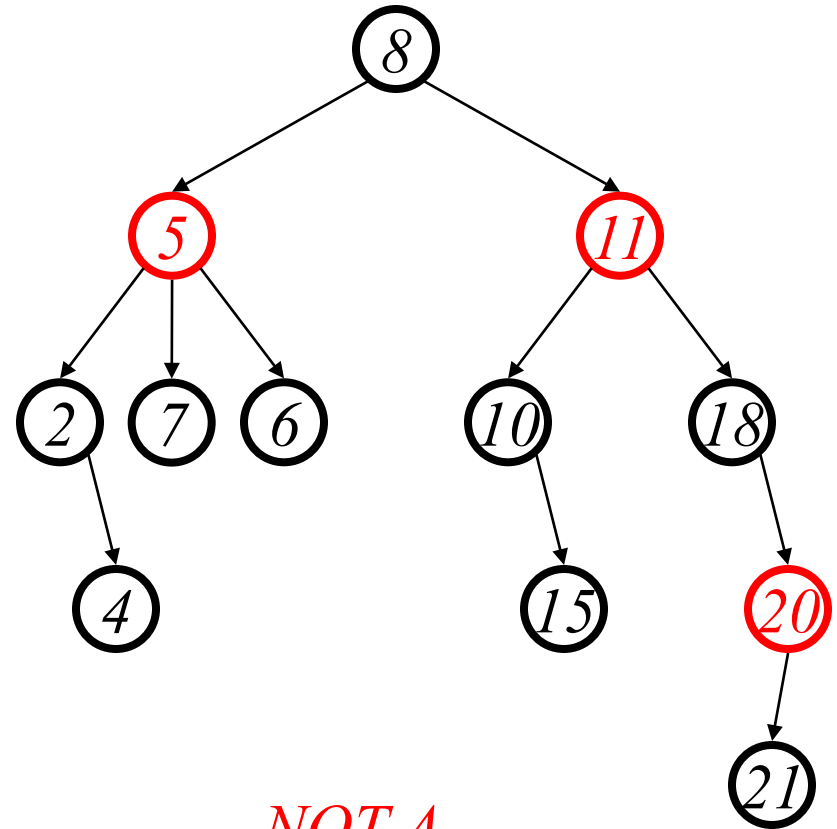
NOT A
BINARY SEARCH TREE

Clicker Question

- The tree on the right is...
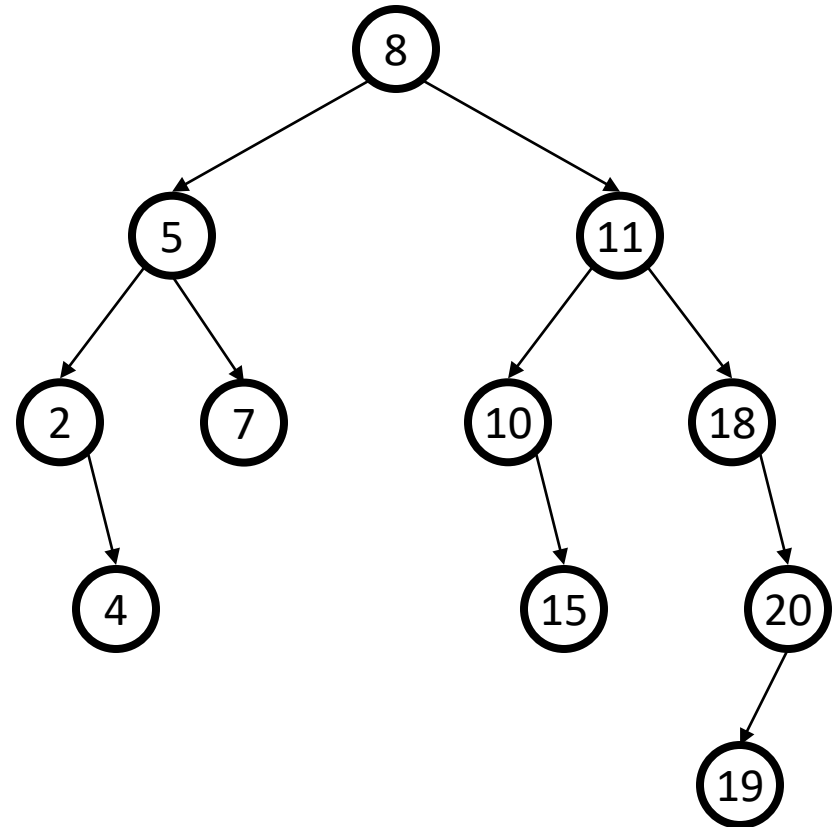
A. a Binary Search Tree.

B. not a Binary Search Tree.

C. hmm... I don't know.

Clicker Question

- The tree on the right is…
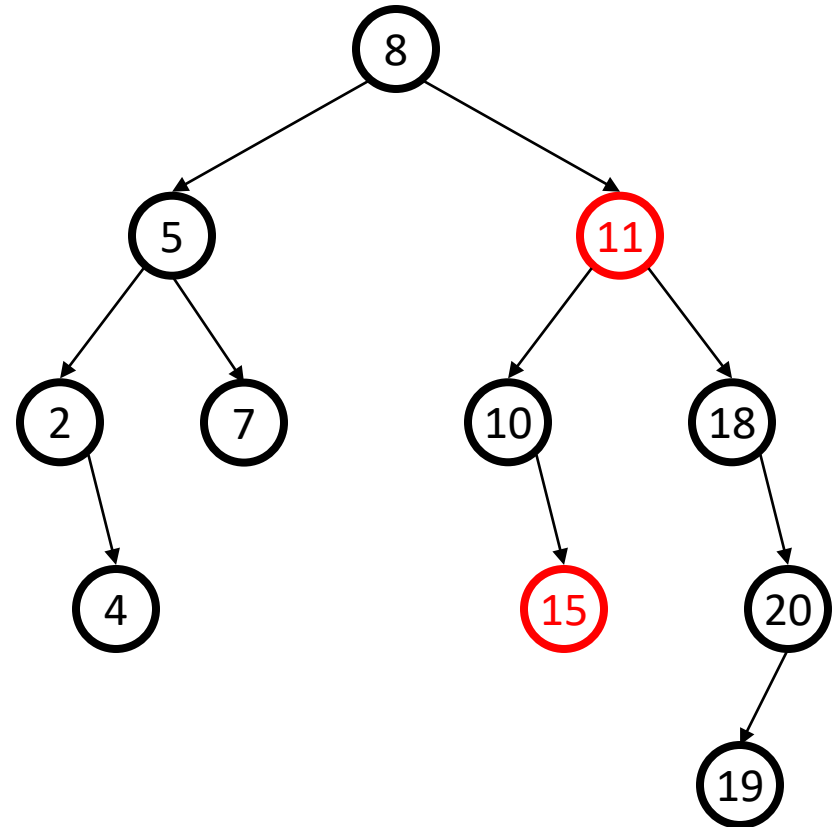
A. a Binary Search Tree.

B. not a Binary Search Tree.

C. hmm... I don't know.

Representing Binary Search Trees


```
class Node2(object):

    def __init__(self, value):
        self.value = value
        self.left = None
        self.right = None
```

Data not shown in figure
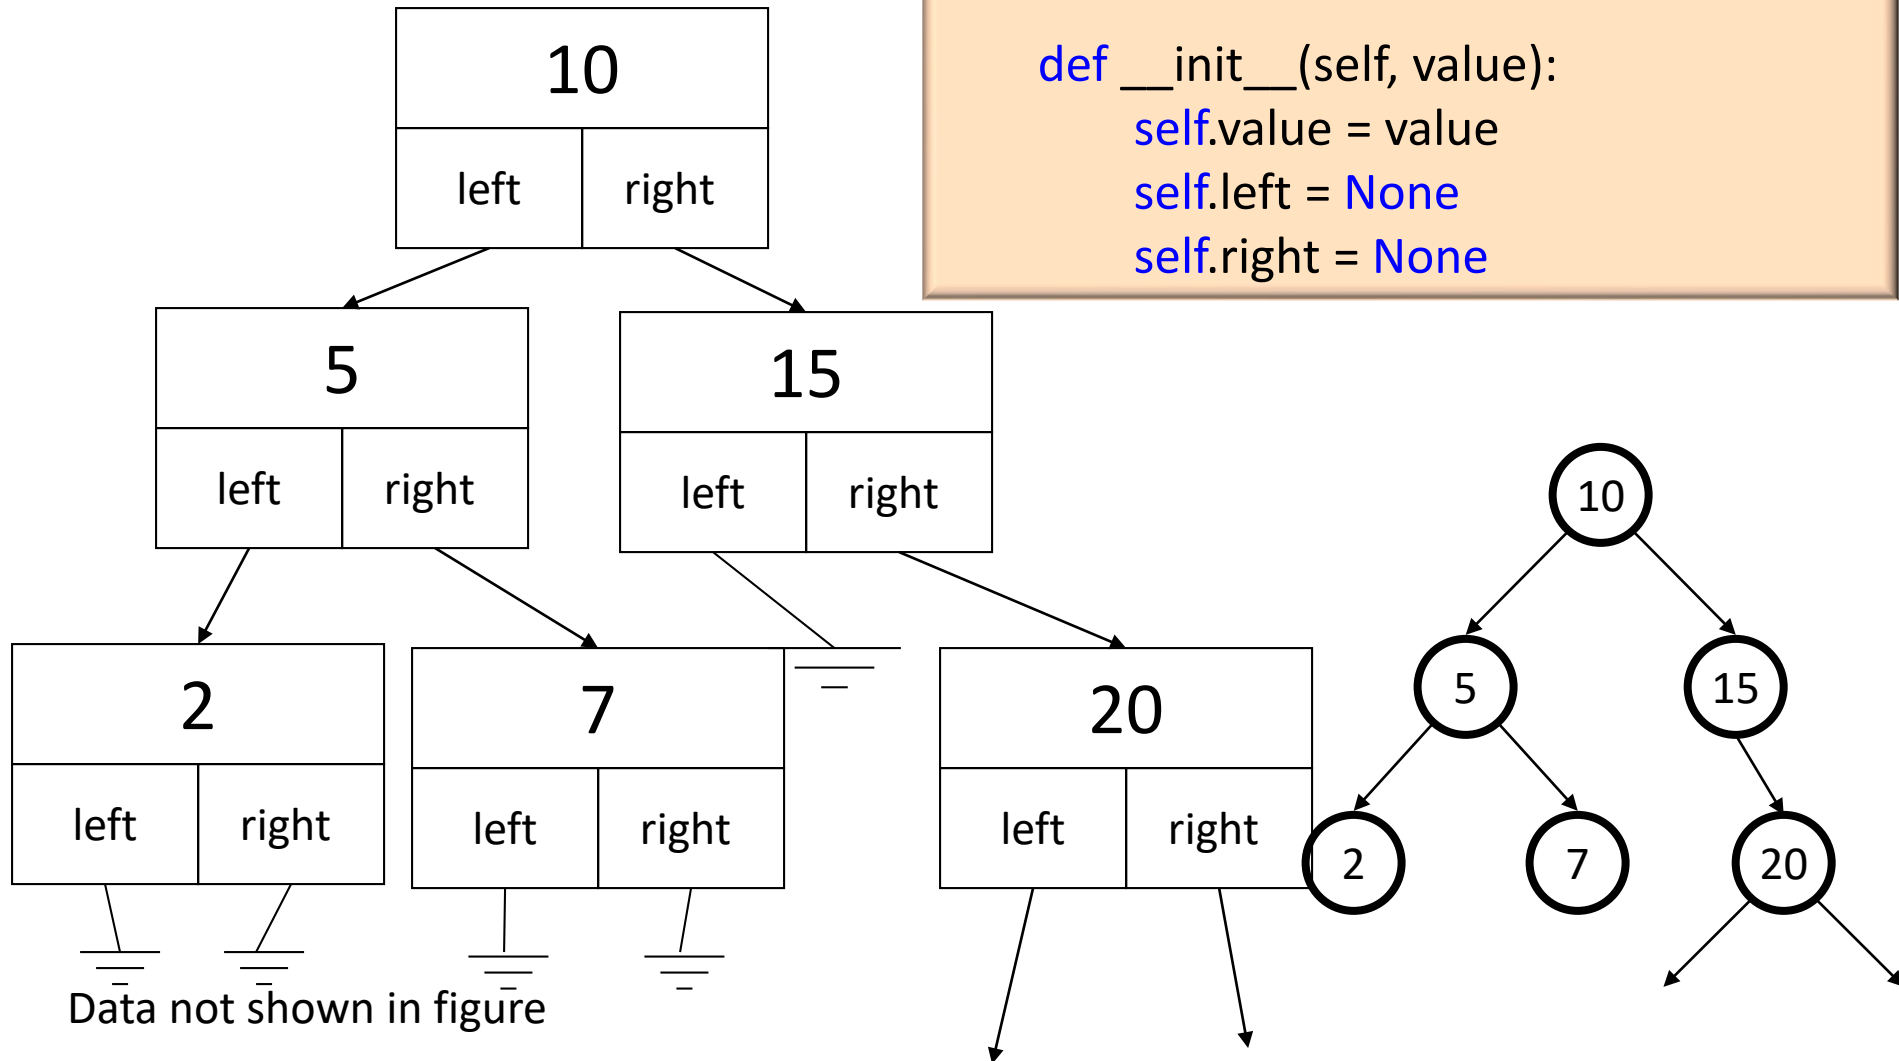
Finding a Node

```python
def search(root,key):
    # Base Cases: root is null or key is present at root
    if root is None or root.value == key:
        return root
    # Key is greater than root's key
    if root.value < key:
        return search(root.right,key)
    # Key is smaller than root's key
    return search(root.left,key)
```
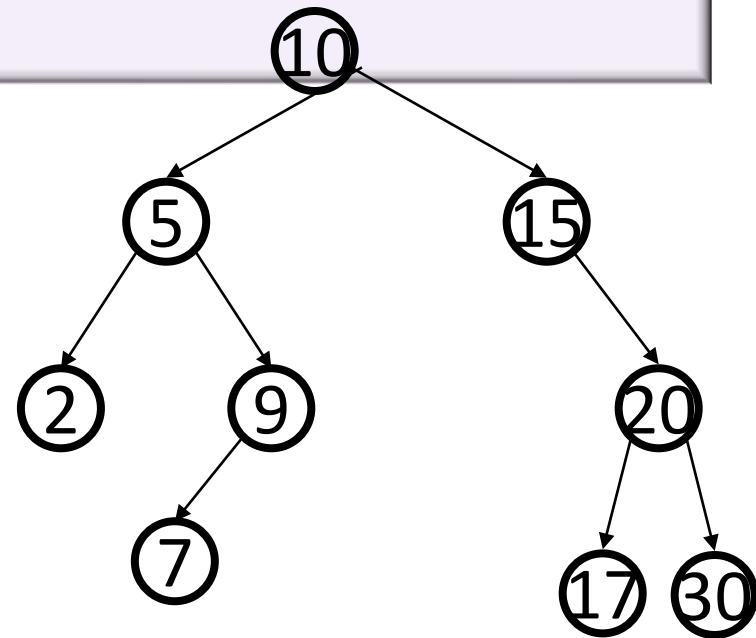
What Makes a Balanced BST Efficient for Searching?
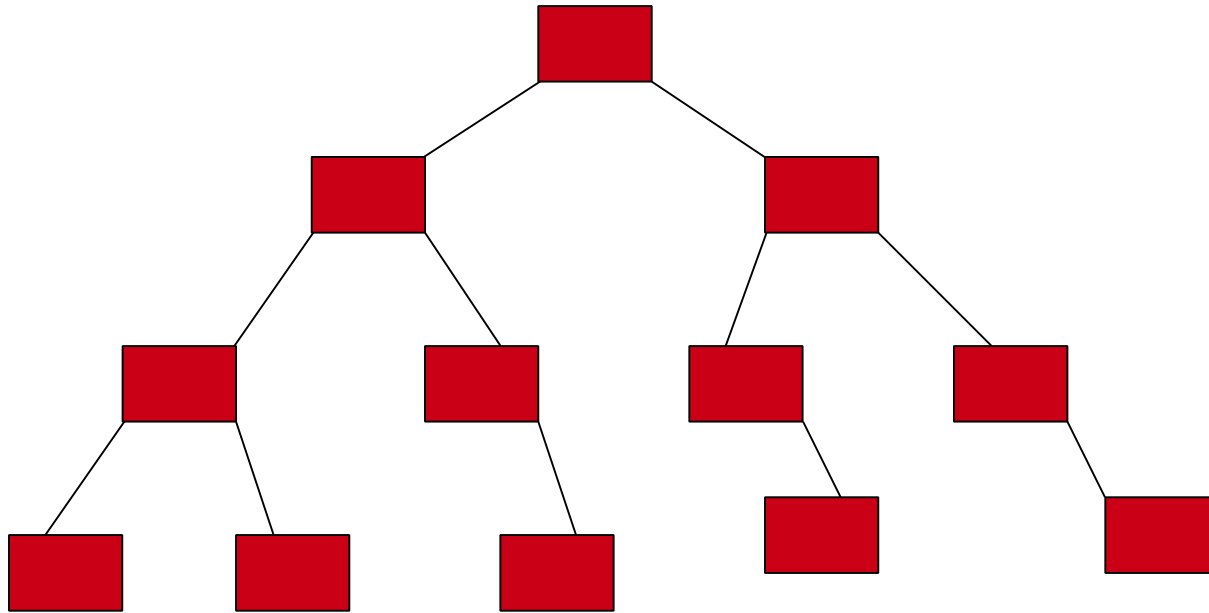
Each step we take as we search the tree reduces the remaining search space by half.

Sample Search

Step 1 →

Each step we take as we search the tree reduces the remaining search space by half.

Sample Search

Step 2

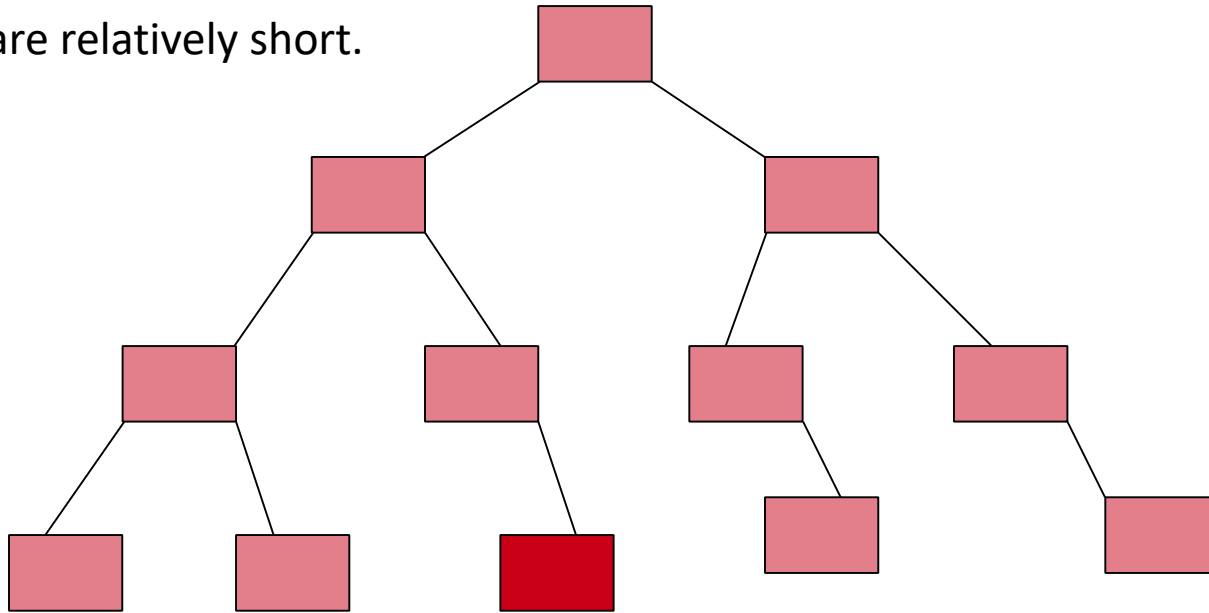Each step we take as we search the tree reduces the remaining search space by half.

Sample Search

Step 3

Each step we take as we search the tree reduces the remaining search space by half.

Sample Search

Step 4
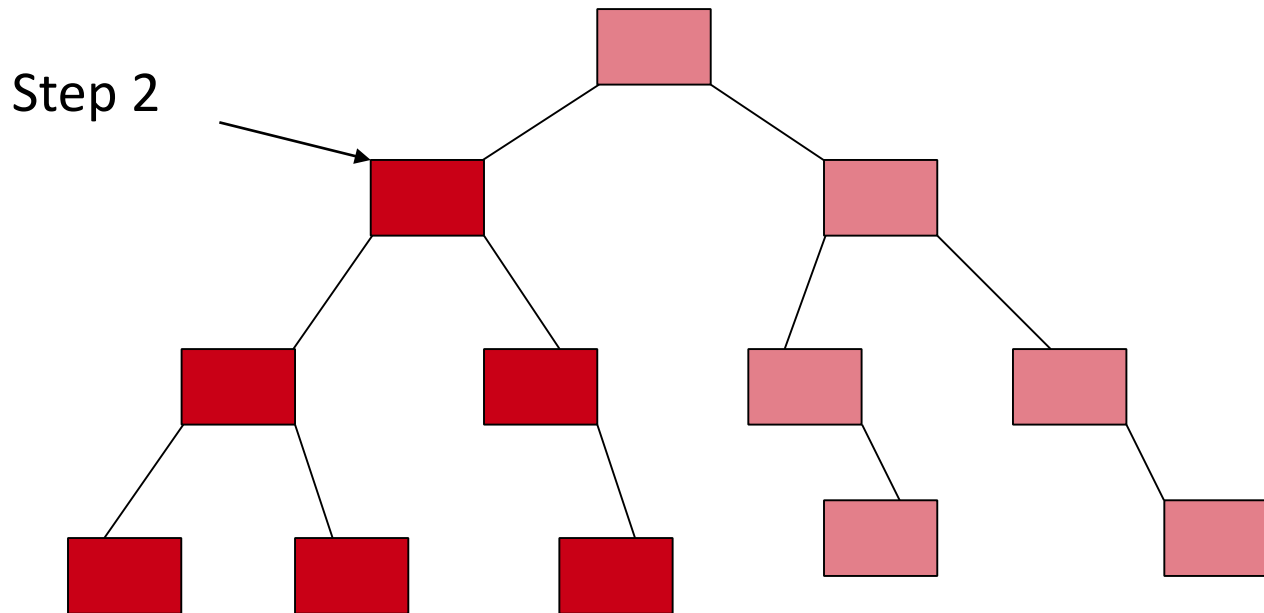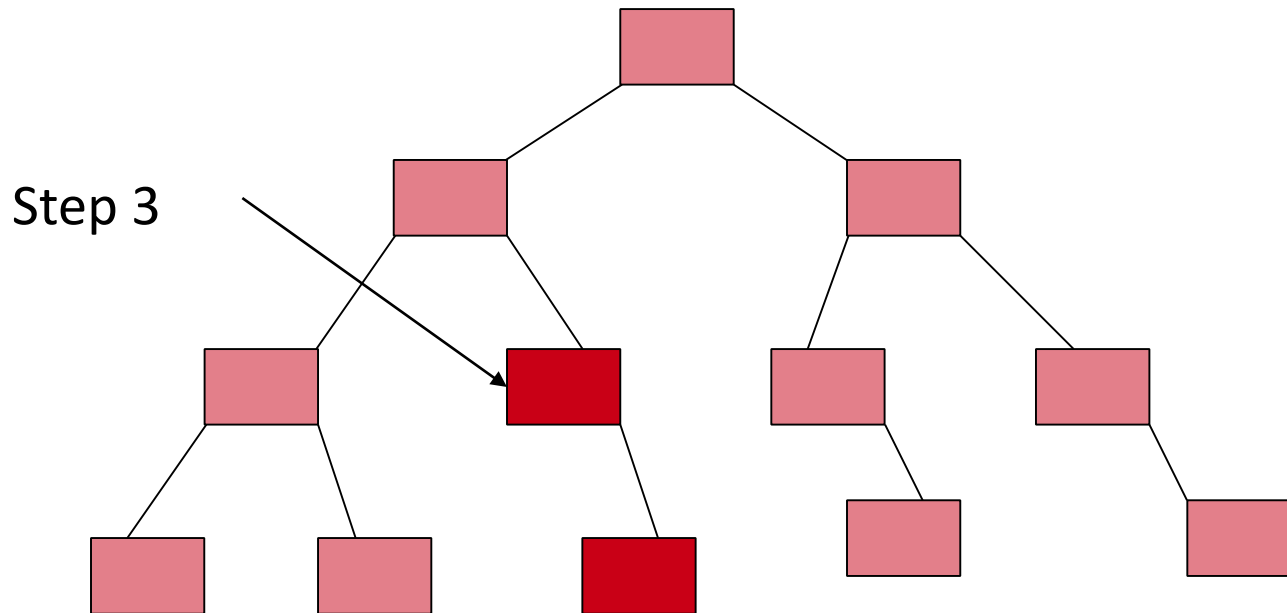
Each step we take as we search the tree reduces the remaining search space by half.

Sample Search

The tree has low height and all
paths from the root node to
other nodes are relatively short.

Each step we take as we search the tree reduces
the remaining search space by half.

Clicker Question

- What is the running time of the find function in BSTs? (height represents the height of the tree)

A. $\Theta(n)$

B. $\Theta(lg\ n)$

C. $\Theta(height)$

D. $\Theta(n^2)$

E. *None of the above*

Clicker Question

- What is the running time of the find function in BSTs? (height represents the height of the tree)

A. $\Theta(n)$

B. $\Theta(lg\ n)$

C. $\Theta(height)$

D. $\Theta(n^2)$

E. None of the above

Unbalanced Trees
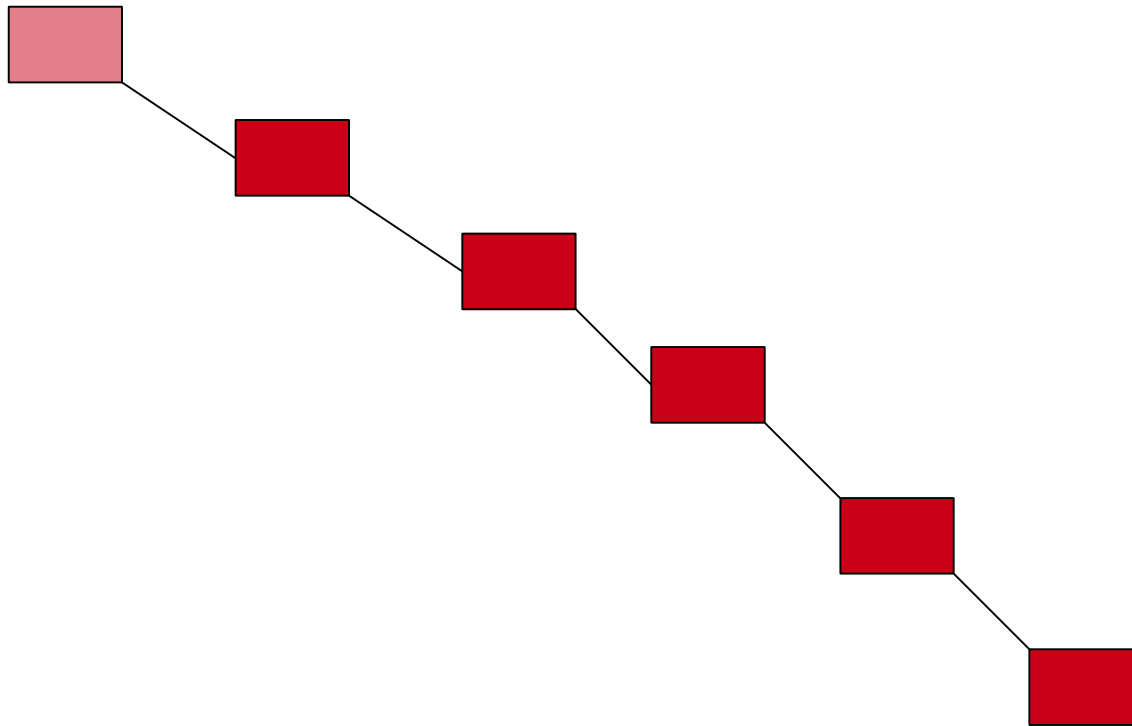
In contrast, this unbalanced tree has long paths from the root to other nodes. It essentially has degenerated to a linked list, which is very slow to search through. Now, with each step we take, we have only reduced the search space by one node.

Time of Search

- Time of search is proportional to the height of the tree

O(lg n)

O(n)
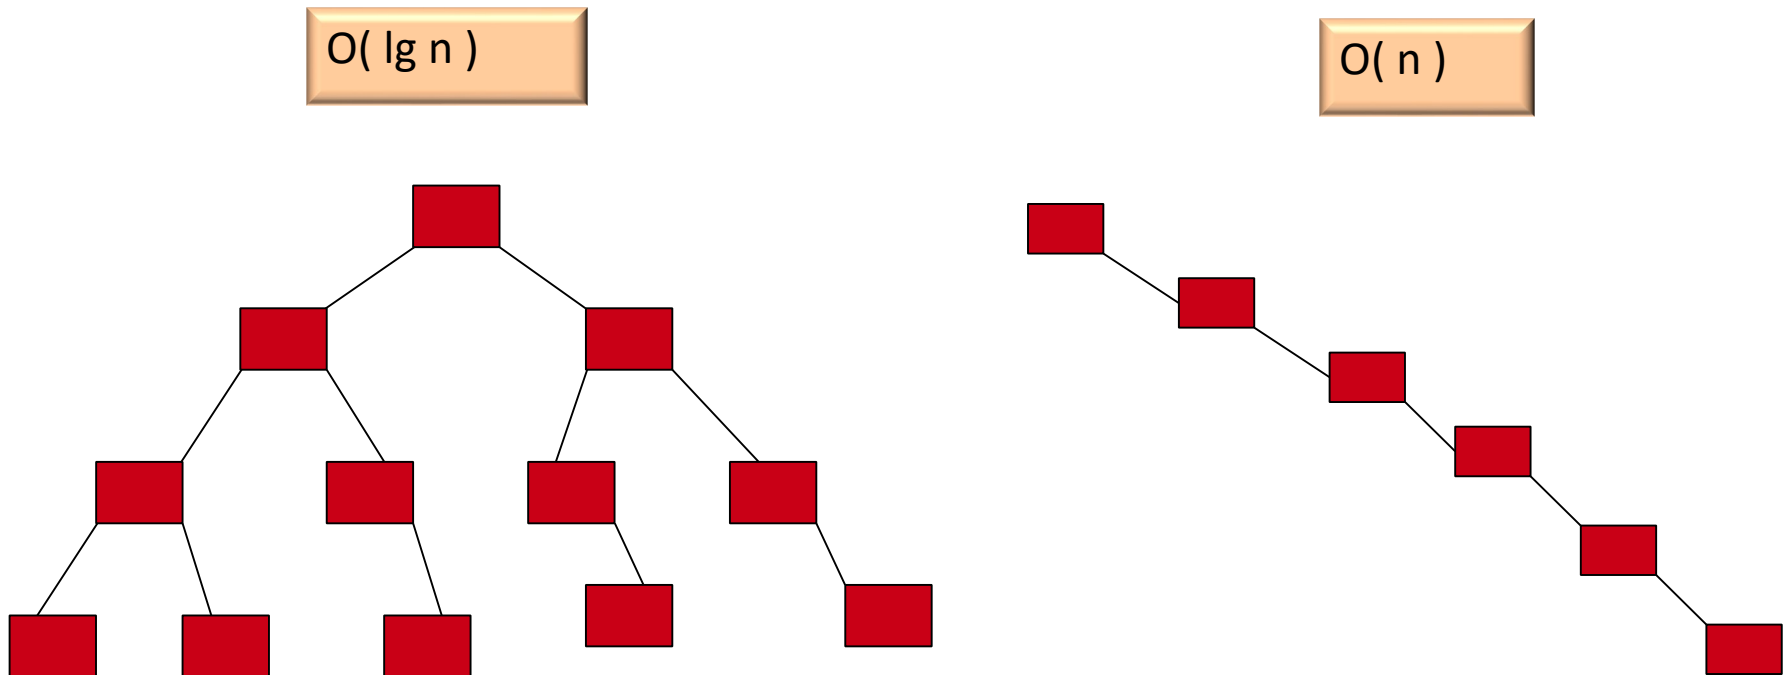
Time of Search

- What does this tell you about strengths and weaknesses of BSTs?

- Using BSTs is only efficient if they are fairly balanced. Whether BSTs are balanced is highly dependent on the order of the values being added.

- There are extensions and improvements to Binary Search Trees such as **AVL trees, B+ trees** or **red-black trees**. I encourage you to read about them, but they are outside of the scope of this course.

Bonus: FindMin/FindMax

- Find minimum

```
def find_min(root):
    current_node = root
    while current_node.left:
        current_node =current_node.left
    return current_node
```

- Find maximum

```
def find_max(root):
    current_node = root
    while current_node.right:
        current_node = current_node.right
    return current_node
```
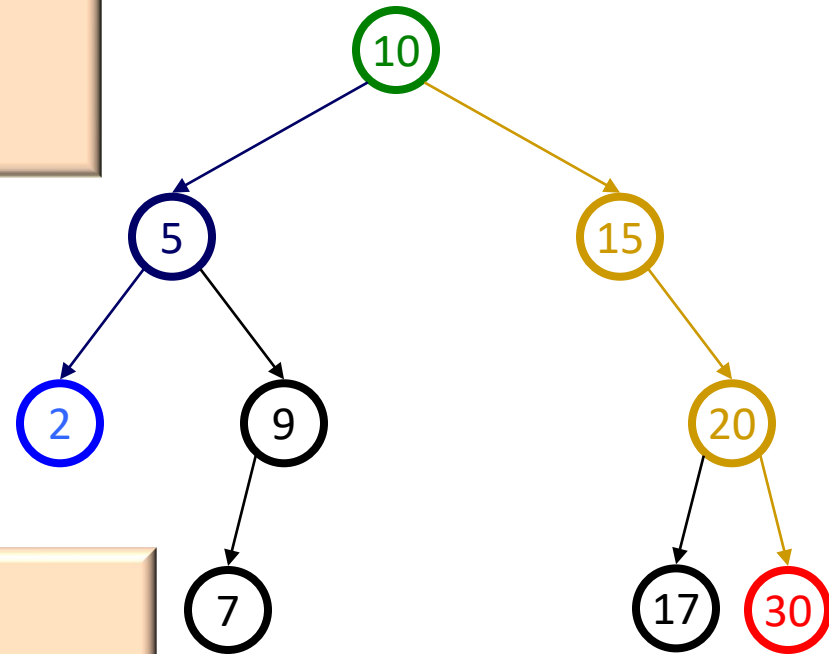
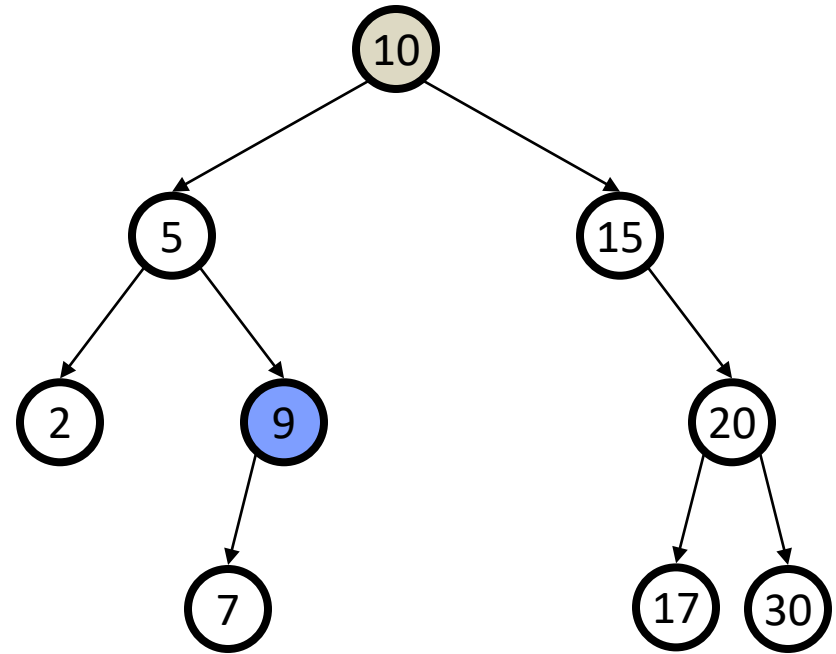# Double Bonus: Successor

Find the next larger node
in a node's subtree.

```
def succ (x):
    return find_min(x.right)
```
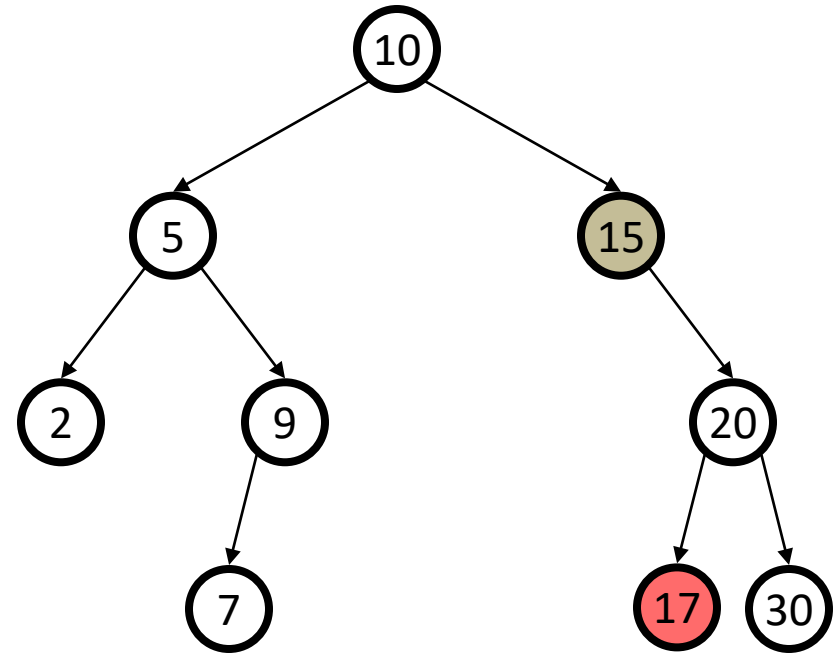
More Double Bonus: Predecessor

Find the next smaller node

in a node's subtree.

```
def Pred (x):
    return find_max(x.left)
```
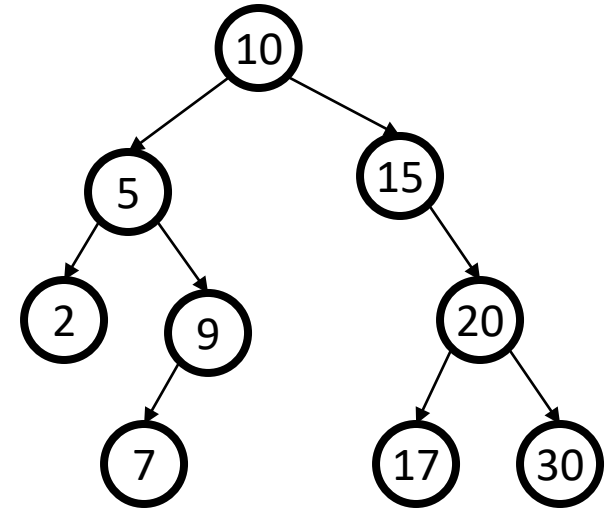
Dictionary ADT

Tree Terminology

Binary Search Trees (BSTs)

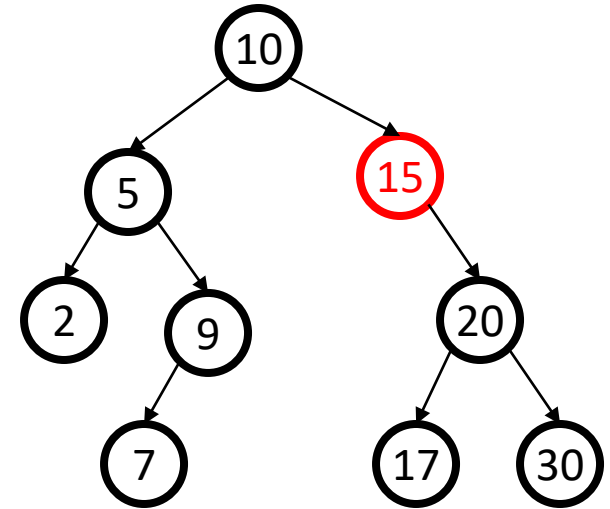**Insertion and Deletion in BSTs**

Insertion Algorithm

1. perform search for value X

2. if X is in tree

    return

    // else search will end at a node. Call it Y

3. if X < Y

    insert X as new left subtree for Y
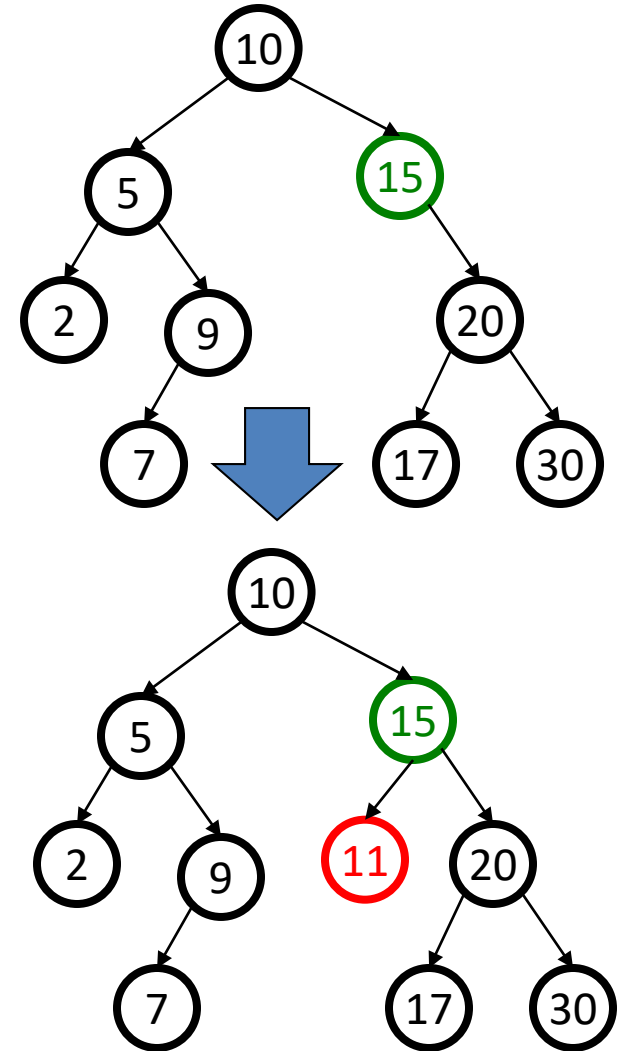
4. if X > Y

    insert X as new right subtree for Y
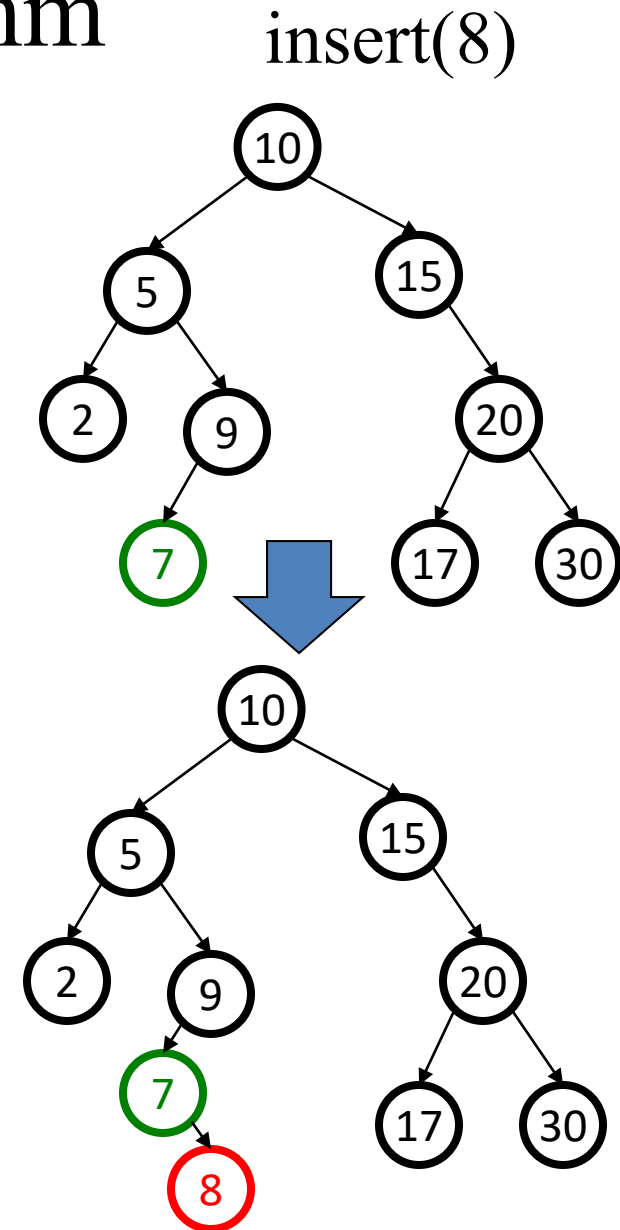
Insertion Algorithm

1. perform search for value X

2. if X is in tree

   return

   // else search will end at a node. Call it Y

3. if X < Y

   insert X as new left subtree for Y

4. if X > Y

   insert X as new right subtree for Y

insert(15)

Insertion Algorithm

1. perform search for value X

2. if X is in tree

    return

    // else search will end at a node. Call it Y

3. if X < Y

    insert X as new left subtree for Y

4. if X > Y

    insert X as new right subtree for Y

insert(11)

Insertion Algorithm

insert(8)

1. perform search for value X

2. if X is in tree

   return

   // else search will end at a node. Call it Y

3. if X < Y

   insert X as new left subtree for Y

4. if X > Y

   insert X as new right subtree for Y

Runtime?

Insertion Algorithm

```python
def insert(root,node):
    if root is None:
        root = node
    else:
        if root.value == node.value: # already exists
            return
        # To be inserted on the right side of current root
        elif root.value < node.value:
            if root.right is None:
                root.right = node
            else:
                insert(root.right, node)
        # To be inserted on the left side of current root
        else:
            if root.left is None:
                root.left = node
            else:
                insert(root.left, node)
```

BuildTree for BSTs

- Suppose the data 1, 2, 3, 4, 5, 6, 7, 8, 9 is inserted into an initially empty BST:
  - in order (1, 2, 3, 4, 5, 6, 7, 8, 9)

  - in reverse order (9, 8, 7, 6, 5, 4, 3, 2, 1)

  - median first, then left median, right median, etc. so: 5, 3, 8, 2, 4, 7, 9, 1, 6

Analysis of BuildTree

- Worst case: $O(n^2)$ as we've seen

- Average case assuming all orderings equally likely turns out to be $O(n \lg n)$.
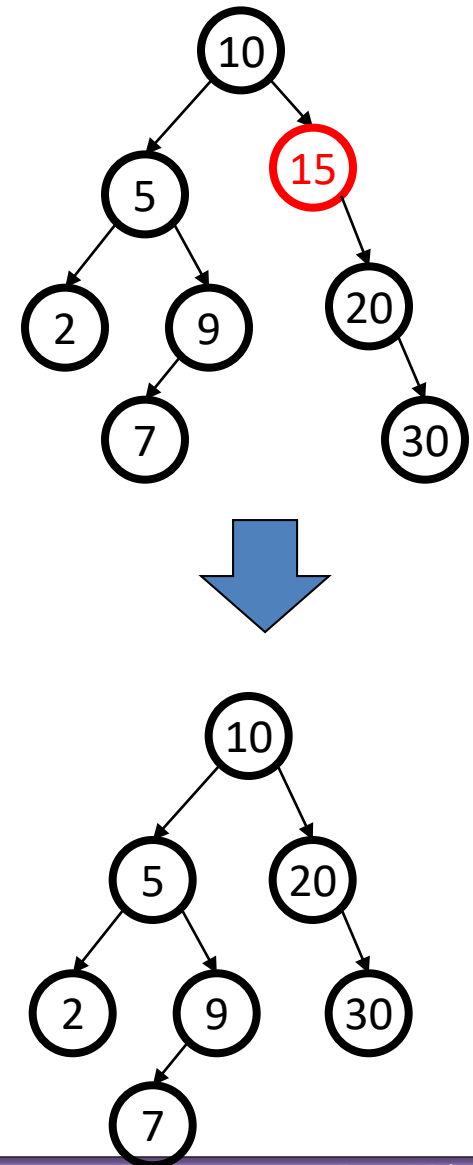
Deletion
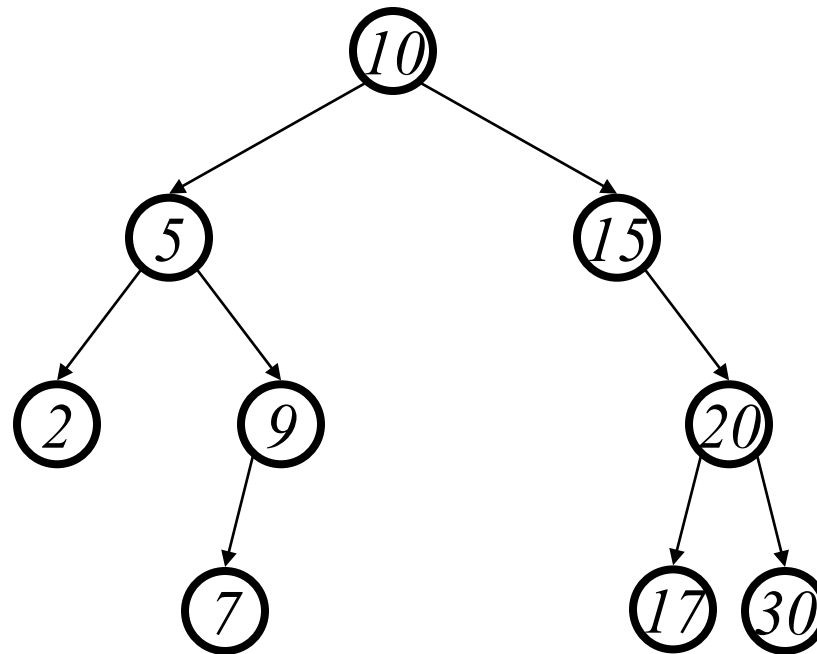
Why might deletion be harder than insertion?

Deletion Algorithm
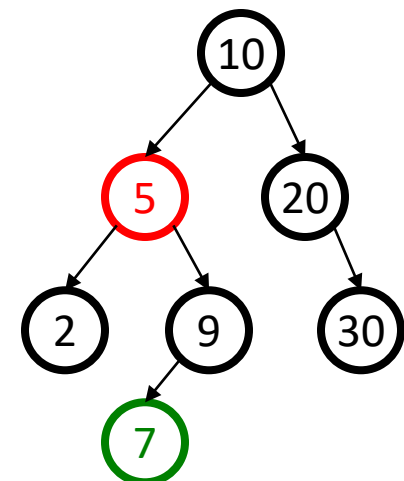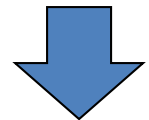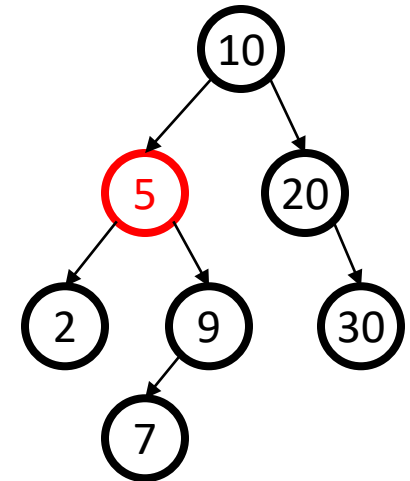
1. perform search for value X

2. if X is not in tree

   return

   //else X is in tree

3. if X has no children,

   delete X

4. else if X has only one child

   promote the unique child to X's place.

   delete X

5. else // X must have two children

   identify X's successor. Call it Y

   replace X with Y

   delete Y // Y $\leq$ 1 children

Deletion Algorithm

1. perform search for value X
2. if X is not in tree
   return

   //else X is in tree

3. if X has no children,

   delete X

4. else if X has only one child

   promote the unique child to X's place.

   delete X

5. else // X must have two children

   identify X's successor. Call it Y

   replace X with Y

   delete Y // Y ≤ 1 children
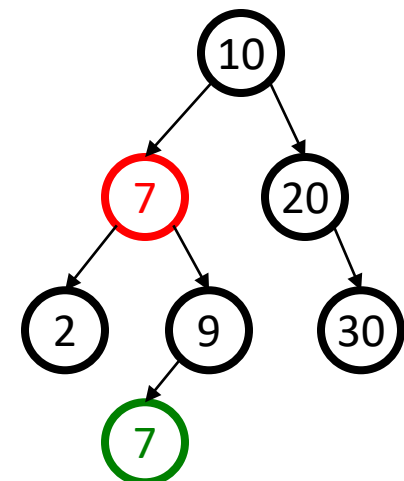
Delete(17)

Deletion Algorithm

Delete(15)

1. perform search for value X

2. if X is not in tree

   return

   //else X is in tree

3. if X has no children,

   delete X

4. else if X has only one child

   promote the unique child to X's place.

   delete X

5. else // X must have two children

   identify X's successor. Call it Y
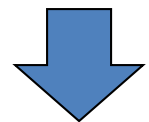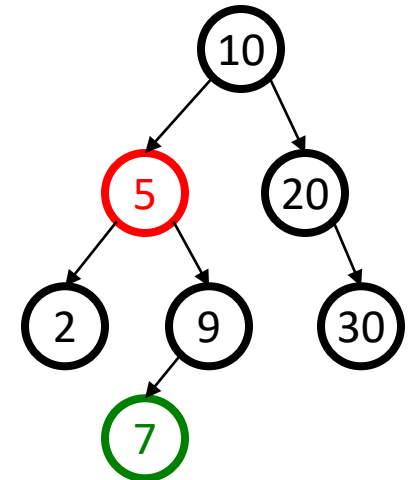
   replace X with Y

   delete Y // Y ≤ 1 children

Deletion Algorithm

1. perform search for value X

2. if X is not in tree

   return

   //else X is in tree

3. if X has no children,

   delete X

4. else if X has only one child

   promote the unique child to X's place.

   delete X

5. else // X must have two children

   identify X's successor. Call it Y
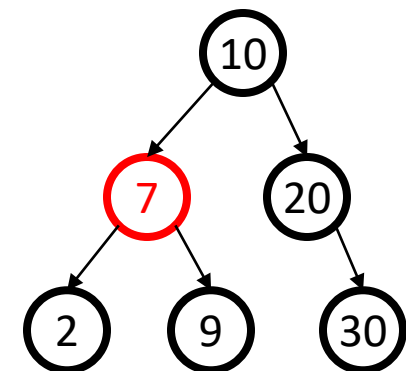
   replace X with Y

   delete Y // Y ≤ 1 children
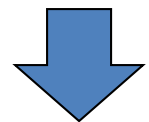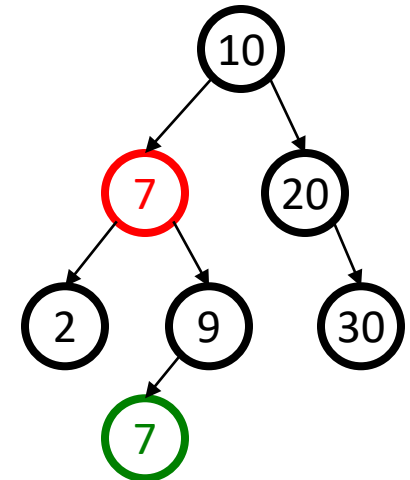
Delete(5)

Deletion Algorithm

1. perform search for value X

2. if X is not in tree

   return

   //else X is in tree

3. if X has no children,

   delete X

4. else if X has only one child

   promote the unique child to X's place.

   delete X

5. else // X must have two children

   identify X's successor. Call it Y

   replace X with Y

   delete Y // Y ≤ 1 children

Delete(5)

Deletion Algorithm

1. perform search for value X

2. if X is not in tree

     return

   //else X is in tree

3. if X has no children,

      delete X

4. else if X has only one child

      promote the unique child to X's place.

      delete X

5. else // X must have two children

      identify X's successor. Call it Y

      replace X with Y

      delete Y // Y ≤ 1 children
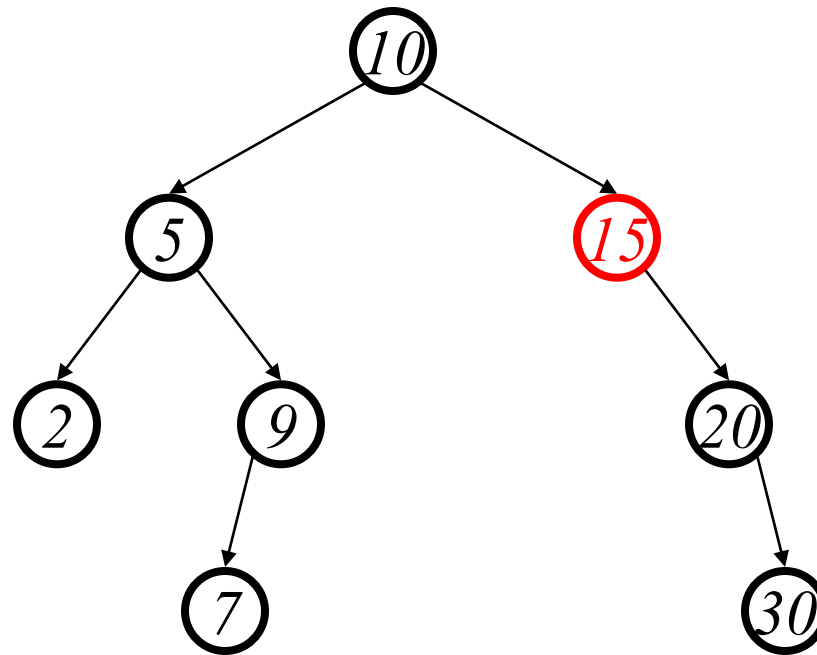
Deletion - Leaf Case
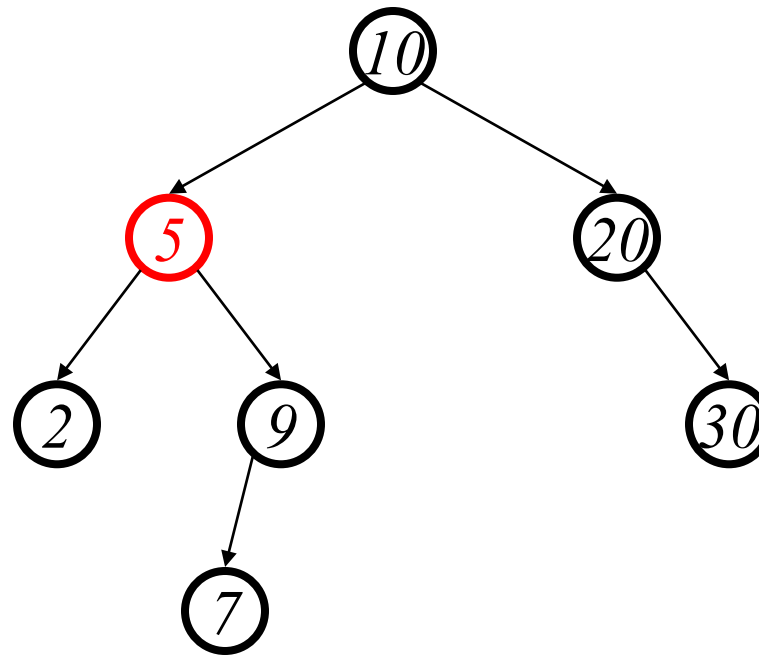
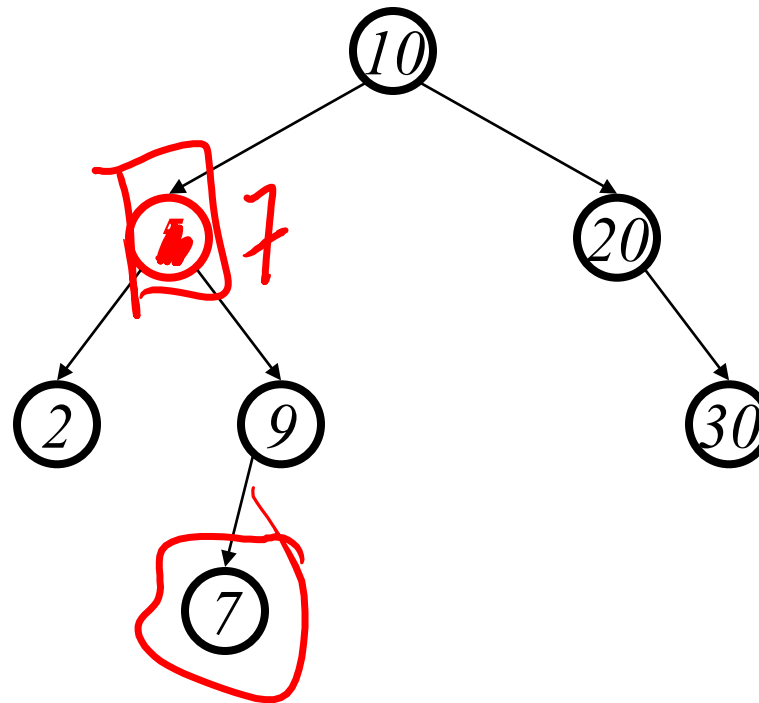*Delete(17)*
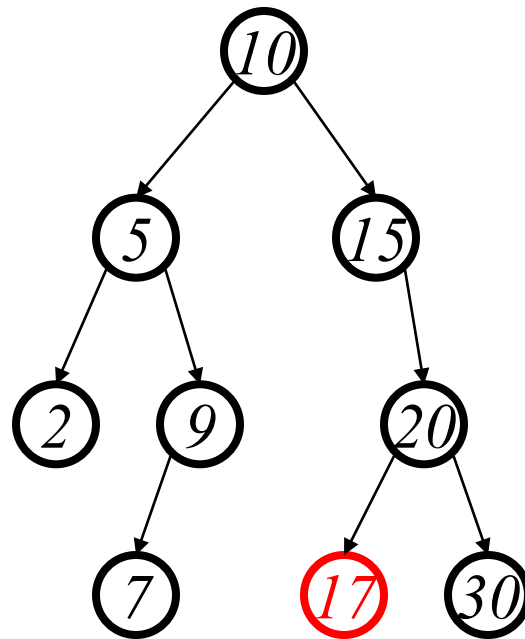
Deletion - One Child Case

Delete(15)

Deletion - Two Child Case
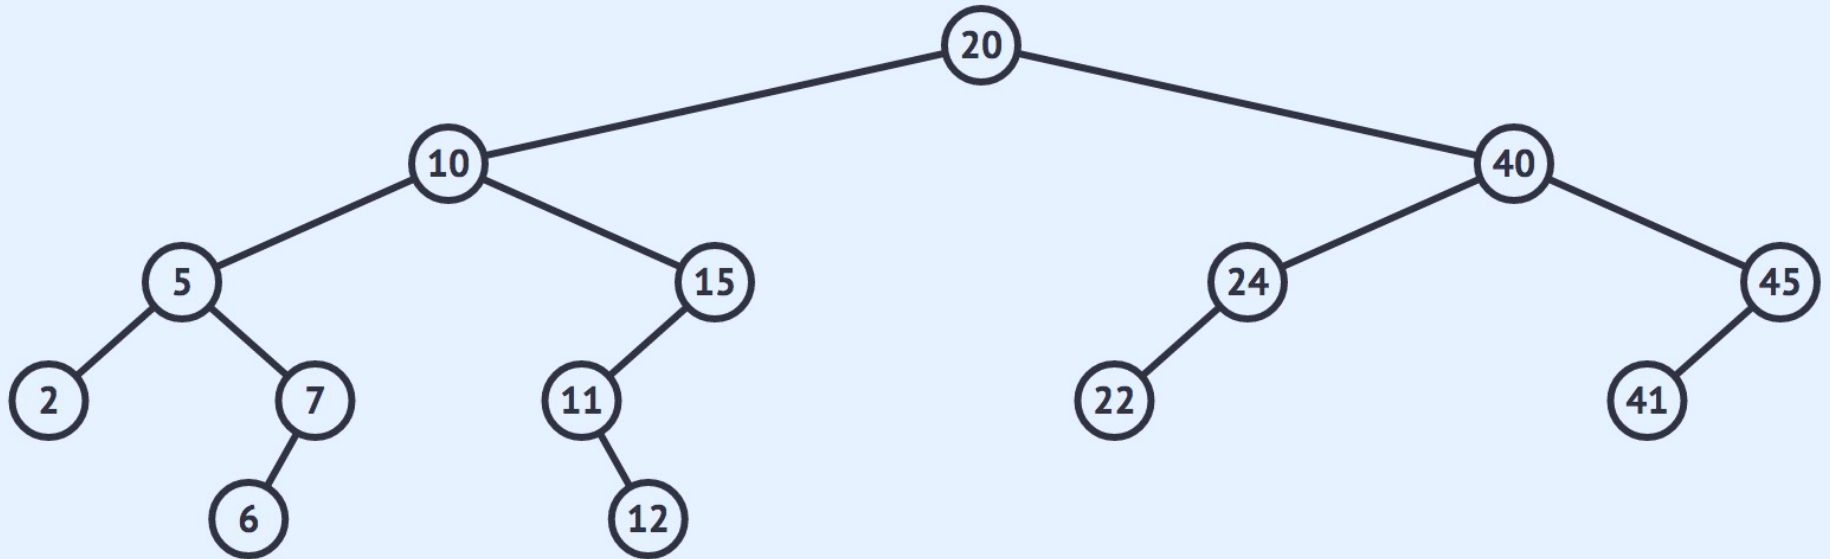
Delete(5)
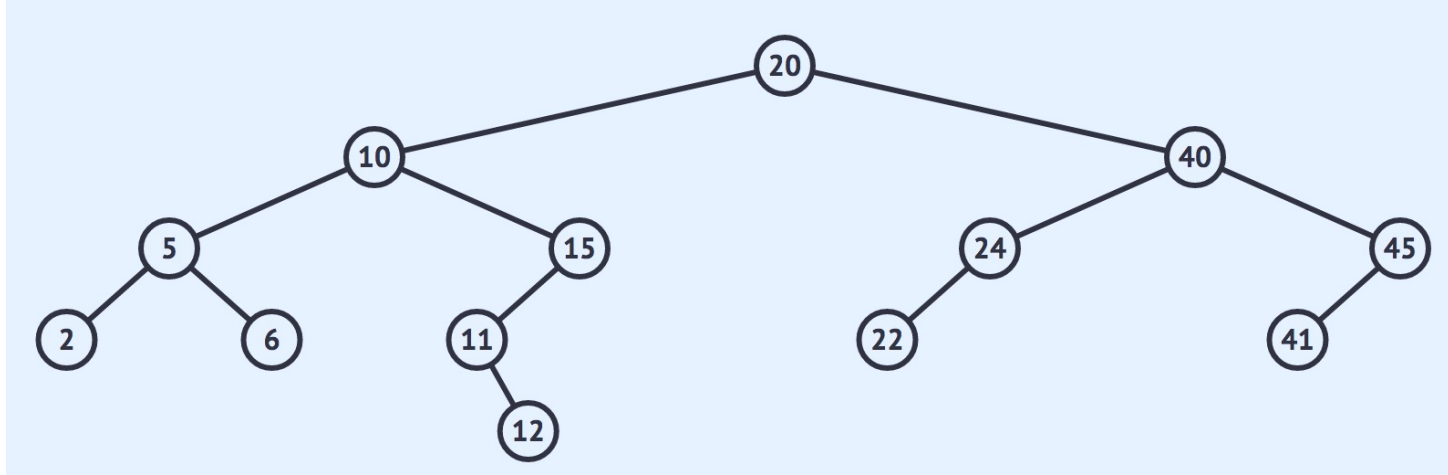
Deletion - Two Child Case

Delete(5)

In-Class Exercise

- Draw the binary search tree, which results from adding the following keys in the given order:
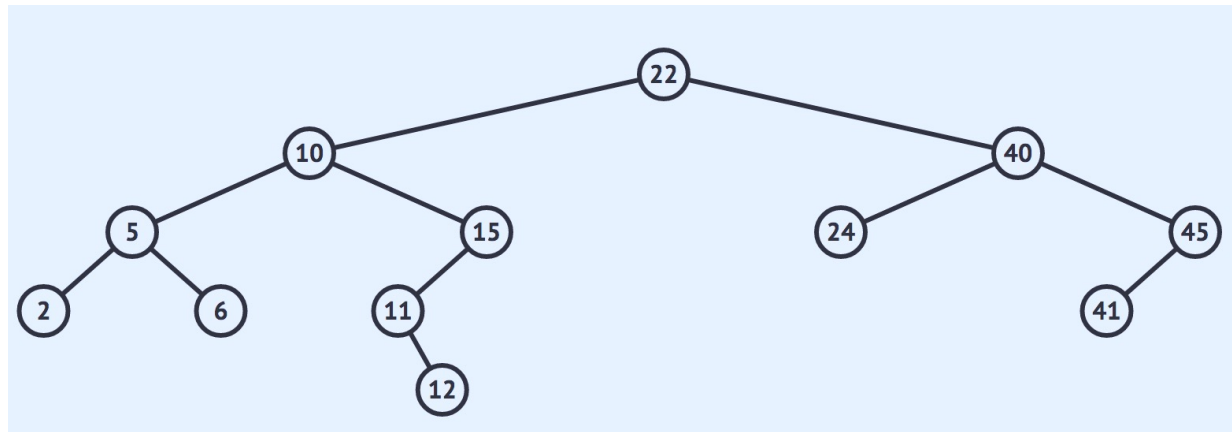 - 20, 10, 40, 5, 7, 2, 15, 11, 12, 6, 24, 22, 45, 41

In-Class Exercise

- From your tree remove key with value 7

- From your tree remove key with value 20 (using successor)

Learning Outcomes

Description	Tag
Describe the different operations that are supported by the dictionary ADT.	Dict-ADT
Demonstrate how arrays can be used for implementation of a dictionary ADT.	Dict-Array
Demonstrate how linked lists can be used for implementation of a dictionary ADT.	Dict-Linkedlist
Describe the properties of binary trees.	Dict-Trees
Describe the properties of Binary Search Trees (BSTs).	
Represent BSTs using linked data structures.	
Insert elements into a BST.	
Delete elements from a BST.	
search for an element in a BST.	
Determine if a given tree is an instance of a BST.	
Compare and contrast algorithms for storing and subsequently retrieving values from a dictionary ADT.	
Reflect on how BSTs and their extensions can be used in DBMS.	